

# Function Calling Conventions in C (x86 Architecture)

## 1. Introduction

Function calling conventions define how function arguments are passed, how the return value is handled, and how the stack is managed during a function call. This varies between 32-bit and 64-bit x86 architectures.

---

## 2. Calling Conventions in x86 (32-bit)

In 32-bit x86, the function calling conventions are mainly:

### 2.1 cdecl (C Declaration)

- Arguments are pushed onto the stack **from right to left**.
- Caller cleans the stack after the function returns.
- Return value is stored in **EAX**.

**Example:**

```
#include <stdio.h>

int sum(int a, int b) {
    return a + b;
}

int main() {
    int result = sum(3, 4);
    printf("Result: %d\n", result);
    return 0;
}
```

**Assembly Code (gcc -m32 -S output snippet):**

```
push    4    ; Push second argument
push    3    ; Push first argument
call    sum   ; Call function
add    esp, 8 ; Clean up the stack (caller cleans up)
```

---

### 2.2 stdcall (Standard Call)

- Arguments are pushed **right to left**.
- **Callee cleans the stack**.
- Used in Windows API functions.

### Assembly Code (MSVC stdcall function):

```
push    4    ; Push second argument
push    3    ; Push first argument
call    sum  ; Call function
; No stack cleanup needed by caller
```

---

## 2.3 fastcall

- First two arguments passed via **ECX, EDX** registers.
- Remaining arguments pushed on the stack.
- **Callee cleans the stack.**

### Example:

```
_fastcall int multiply(int a, int b) {
    return a * b;
}
```

---

## 2.4 thiscall (Used for C++ Member Functions)

- **this** pointer is passed in **ECX**.
- Other arguments are passed on the stack.

## 3. Calling Conventions in x86-64 (64-bit)

In 64-bit architecture, function arguments are primarily passed in registers, reducing the reliance on the stack, leading to better performance.

### 3.1 System V ABI (Linux/macOS)

- First **six** integer/pointer arguments are passed in **RDI, RSI, RDX, RCX, R8, R9**.
- Floating-point arguments are passed in **XMM0–XMM7**.
- Remaining arguments are passed on the stack.
- Caller cleans the stack.
- The return value is stored in **RAX** (or **XMM0** for floating-point values).

### Assembly Code Example:

```
mov    edi, 3    ; First argument (RDI)
mov    esi, 4    ; Second argument (RSI)
call   sum      ; Call function
```

---

### 3.2 Microsoft x64 ABI (Windows)

- First **four** integer arguments are passed in **RCX, RDX, R8, R9**.

- Floating-point arguments are passed in **XMM0–XMM3**.
- Remaining arguments are passed on the stack.
- Caller cleans the stack.
- The return value is stored in **RAX** (or XMM0 for floating-point values).
- A "shadow space" (32 bytes) is allocated on the stack before function calls for register arguments.

### Example Function Call in Windows x64 ABI:

```
mov    rcx, 3      ; First argument (RCX)
mov    rdx, 4      ; Second argument (RDX)
call   sum         ; Call function
```

---

## 3.3 Changing the Calling Convention in x86-64

Unlike x86-32, where multiple calling conventions exist, x86-64 architectures have standard ABIs. However, you can override them manually:

### Using `__attribute__((sysv_abi))` (GCC/Clang, Linux/macOS)

```
#include <stdio.h>

int __attribute__((sysv_abi)) sum(int a, int b) {
    return a + b;
}
```

This forces the function to use System V calling convention explicitly.

### Using `__vectorcall` (MSVC, Windows)

```
#include <stdio.h>

int __vectorcall sum(int a, int b) {
    return a + b;
}
```

This modifies how arguments are passed in registers (favoring SIMD registers for floating-point).

### Inline Assembly to Override Conventions

You can manually override register assignments before function calls using inline assembly:

```
mov    rdi, 3      ; Force System V ABI style
mov    rsi, 4
call   sum
```

---

## 3.4 Comparison: System V vs Microsoft x64 ABI

Feature	System V (Linux/macOS)	Microsoft x64 (Windows)
Integer Arguments	RDI, RSI, RDX, RCX, R8, R9	RCX, RDX, R8, R9

Feature	System V (Linux/macOS)	Microsoft x64 (Windows)
Floating Arguments	XMM0–XMM7	XMM0–XMM3
Stack Cleanup	Caller	Caller
Shadow Space	None	32 bytes
Return Value	RAX (XMM0 for FP)	RAX (XMM0 for FP)

---

## 4. Observing Calling Conventions Programmatically

We can inspect calling conventions using **disassembly and debugging**.

### 4.1 Using GDB (Linux/macOS)

Compile with debugging symbols:

```
gcc -m32 -g -o test test.c
```

Run in GDB:

```
gdb ./test
break sum
run
info registers # Check argument registers or stack usage
```

### 4.2 Using objdump

To view the assembly:

```
gcc -m32 -S test.c -o test.s
objdump -d ./test
```

### 4.3 Using MSVC Debugger (Windows)

Compile:

```
cl /FAs /Zi test.c
```

Run in Debug Mode and observe register values.

---

## 5. Conclusion

Understanding function calling conventions helps in debugging, reverse engineering, and low-level optimization. The key differences between **x86 (32-bit)** and **x86-64 (64-bit)** are in how arguments are passed and how stack cleanup is handled.

- **x86-32:** Mostly stack-based argument passing.
- **x86-64:** Register-based argument passing for improved performance.

- You can modify the calling convention using attributes like `sysv_abi`, `vectorcall`, or inline assembly.

## Application Binary Interface (ABI) - Short Notes

An **Application Binary Interface (ABI)** defines the low-level interface between software (programs) and the underlying operating system or hardware. It ensures that binaries compiled with different compilers or languages can interact correctly.

### Key Components of an ABI

1. **Calling Conventions** – Defines how function arguments are passed (registers vs stack), return values, and stack cleanup responsibilities.
2. **Register Usage** – Specifies which registers are preserved across function calls and which are caller/callee-saved.
3. **Stack Frame Layout** – Determines how function stack frames are organized, including local variables and return addresses.
4. **Data Alignment & Padding** – Defines how data structures are aligned in memory for performance and compatibility.
5. **Exception Handling** – Specifies how exceptions are propagated and handled across function calls.
6. **System Calls Interface** – Defines how user-space programs interact with the operating system kernel.

### Common ABIs

- **System V ABI (Linux/macOS x86-64)** – Uses registers RDI, RSI, RDX, RCX, R8, R9 for function arguments.
- **Microsoft x64 ABI (Windows)** – Uses RCX, RDX, R8, R9 for function arguments, with shadow space.
- **ARM EABI (Embedded Systems)** – Defines ARM register conventions for efficient execution.

In C, you can implement a **variable argument list function** using the `<stdarg.h>` library, which provides macros to handle functions with a variable number of arguments.

---

## Steps to Implement a Variadic Function

1. **Include `<stdarg.h>`** for handling variable arguments.
  2. **Use `va_list`** to declare a variable to traverse the arguments.
  3. **Initialize the list using `va_start`**, passing the last known fixed parameter.
  4. **Extract each argument using `va_arg`**, specifying the expected type.
  5. **Clean up using `va_end`** to avoid undefined behavior.
- 

## Example: A Function to Calculate the Sum of N Numbers

```
#include <stdarg.h>
#include <stdio.h>

// Function to compute sum of given numbers
int sum(int count, ...) {
    va_list args;           // Declare variable argument list
    va_start(args, count); // Initialize the list, 'count' is the last fixed
    argument

    int total = 0;
    for (int i = 0; i < count; i++) {
        total += va_arg(args, int); // Fetch the next argument
    }

    va_end(args); // Cleanup
    return total;
}

int main() {
    printf("Sum of 3, 4, 5: %d\n", sum(3, 3, 4, 5));
    printf("Sum of 10, 20, 30, 40: %d\n", sum(4, 10, 20, 30, 40));
    return 0;
}
```

### Explanation:

- The first parameter `count` determines how many arguments follow.
  - The `va_list args` holds the arguments.
  - `va_arg(args, int)` retrieves the next argument in the list.
  - `va_end(args)` releases the memory.
-

## Example: A Function That Accepts Different Data Types

```
#include <stdarg.h>
#include <stdio.h>

// Function to print different data types
void printValues(int num, ...) {
    va_list args;
    va_start(args, num);

    for (int i = 0; i < num; i++) {
        switch (i) {
            case 0: printf("Integer: %d\n", va_arg(args, int)); break;
            case 1: printf("Double: %f\n", va_arg(args, double)); break;
            case 2: printf("Char: %c\n", va_arg(args, int)); break; // char is
promoted to int
        }
    }

    va_end(args);
}

int main() {
    printValues(3, 42, 3.14, 'A');
    return 0;
}
```

### Key Notes:

- double arguments are automatically promoted (hence `va_arg(args, double)`).
- char arguments are promoted to int (hence `va_arg(args, int)`).

---

## Limitations of Variadic Functions

- No type safety (incorrect types may lead to undefined behavior).
- No way to determine the number of arguments without an explicit count.
- Cannot directly handle struct types (must use pointers).

## Implementing a `printf`-like Function Using Variadic Arguments in C

The `printf` function in C is a perfect example of a **variadic function** that accepts different argument types based on format specifiers. Below is a simplified implementation of a custom `my_printf` function.

---

### Custom `printf` Implementation

```
#include <stdarg.h>
#include <stdio.h>

// Custom printf function
void my_printf(const char *format, ...) {
    va_list args;
    va_start(args, format);

    for (const char *ptr = format; *ptr != '\0'; ptr++) {
        if (*ptr == '%' && *(ptr + 1) != '\0') {
            ptr++; // Move to the format specifier
            switch (*ptr) {
                case 'd': {
                    int i = va_arg(args, int);
                    printf("%d", i);
                    break;
                }
                case 'f': {
                    double f = va_arg(args, double);
                    printf("%f", f);
                    break;
                }
                case 'c': {
                    char c = (char)va_arg(args, int); // char is promoted to int
                    printf("%c", c);
                    break;
                }
                case 's': {
                    char *s = va_arg(args, char *);
                    printf("%s", s);
                    break;
                }
                default:
                    putchar('%');
                    putchar(*ptr);
            }
        } else {
            putchar(*ptr);
        }
    }

    va_end(args);
}

int main() {
    my_printf("Hello %s! Your score is %d and percentage is %f. Grade: %c\n",
              "Alice", 95, 92.5, 'A');
    return 0;
}
```

}

---

## How It Works:

1. `va_start(args, format)` initializes the argument list.
  2. Iterates through the format string character by character.
  3. When % is found, the next character determines the data type:
    - o `%d` → Extracts an `int` using `va_arg(args, int)`.
    - o `%f` → Extracts a `double` using `va_arg(args, double)`.
    - o `%c` → Extracts a `char` (promoted to `int`).
    - o `%s` → Extracts a `char*` (string).
  4. Calls `printf` for each extracted value.
  5. If no format specifier is found, prints the character as is.
  6. `va_end(args)` is called to clean up.
- 

## Example Output:

Hello Alice! Your score is 95 and percentage is 92.500000. Grade: A

---

## Key Takeaways:

- `va_list`, `va_start`, `va_arg`, and `va_end` are essential for variadic functions.
- `printf`-like functions need to handle multiple data types carefully.
- `char` arguments are promoted to `int`, and `float` is promoted to `double` when passed to a variadic function.