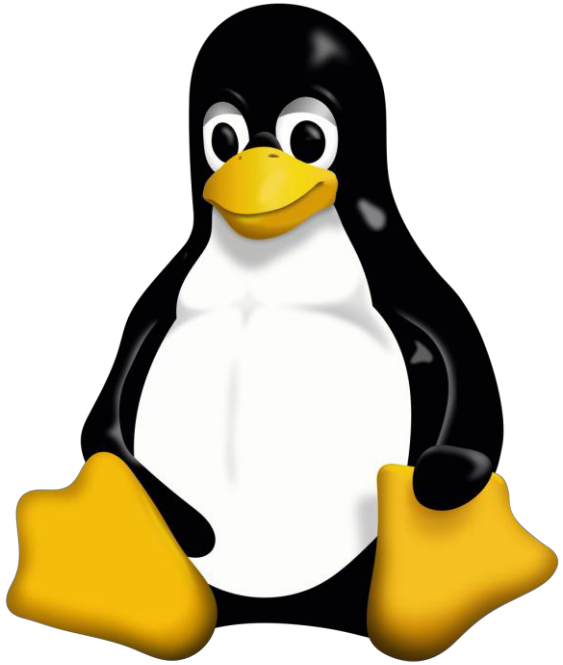

Linux Character Device Driver

Sunbeam Infotech

Register character device

- Each char device is represented by struct cdev.
- *cdev_map* is global array (hash table) to keep track of all char devices.
- cdev_map is object of struct kobj_map.
  - key = device major number
  - hash function = major % 255
  - value = struct probe
    - void *data = struct cdev
- (2) Add device into char device database *cdev_map*.
  - cdev_init(&cdev, &fops);
  - cdev_add(&cdev, devno, dev_count);
- Added device can be removed.
  - cdev_del(&cdev);

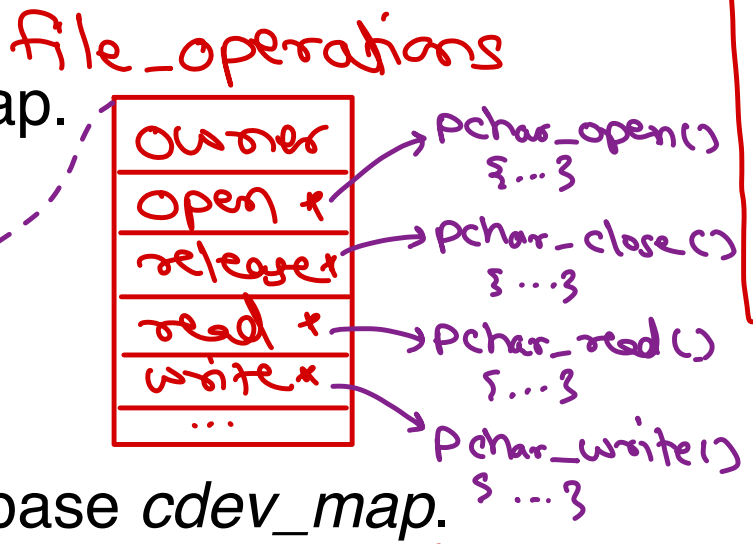


```
struct kobj_map {
  struct probe {
    dev_t dev;
    unsigned long range;
    ...
    void *data;
  } *probes[255];
  ...
};
```

struct kobj_map cdev_map;

```
struct cdev {
    struct kobject kobj;
    struct module *owner;
    struct file_operations *ops;
    struct list_head list;
    dev_t dev;
    unsigned count;
};
```

cdev_add()

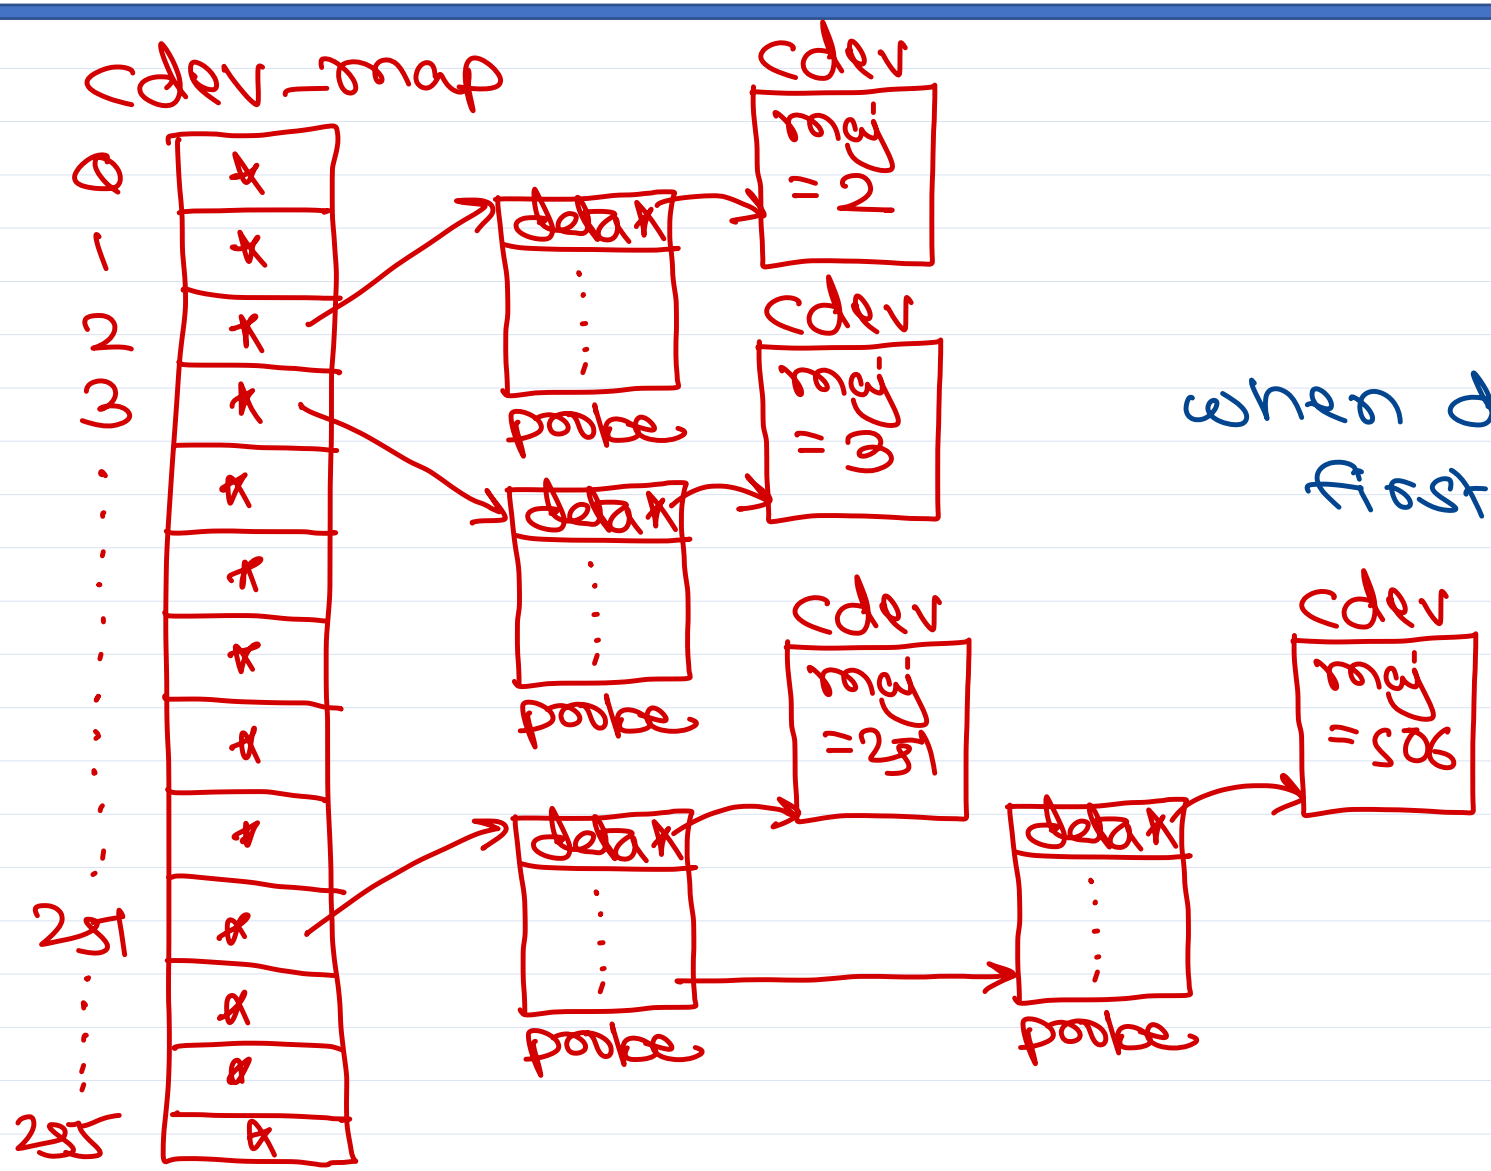

key = major e.g. 506

hash fn = key % 255

e.g. 506 % 255 = 251

when device file is opened for first time, its cdev will be searched from this cdev_map using dev no (i_rdev).

Execution Flow of Pseudo Char Device Driver

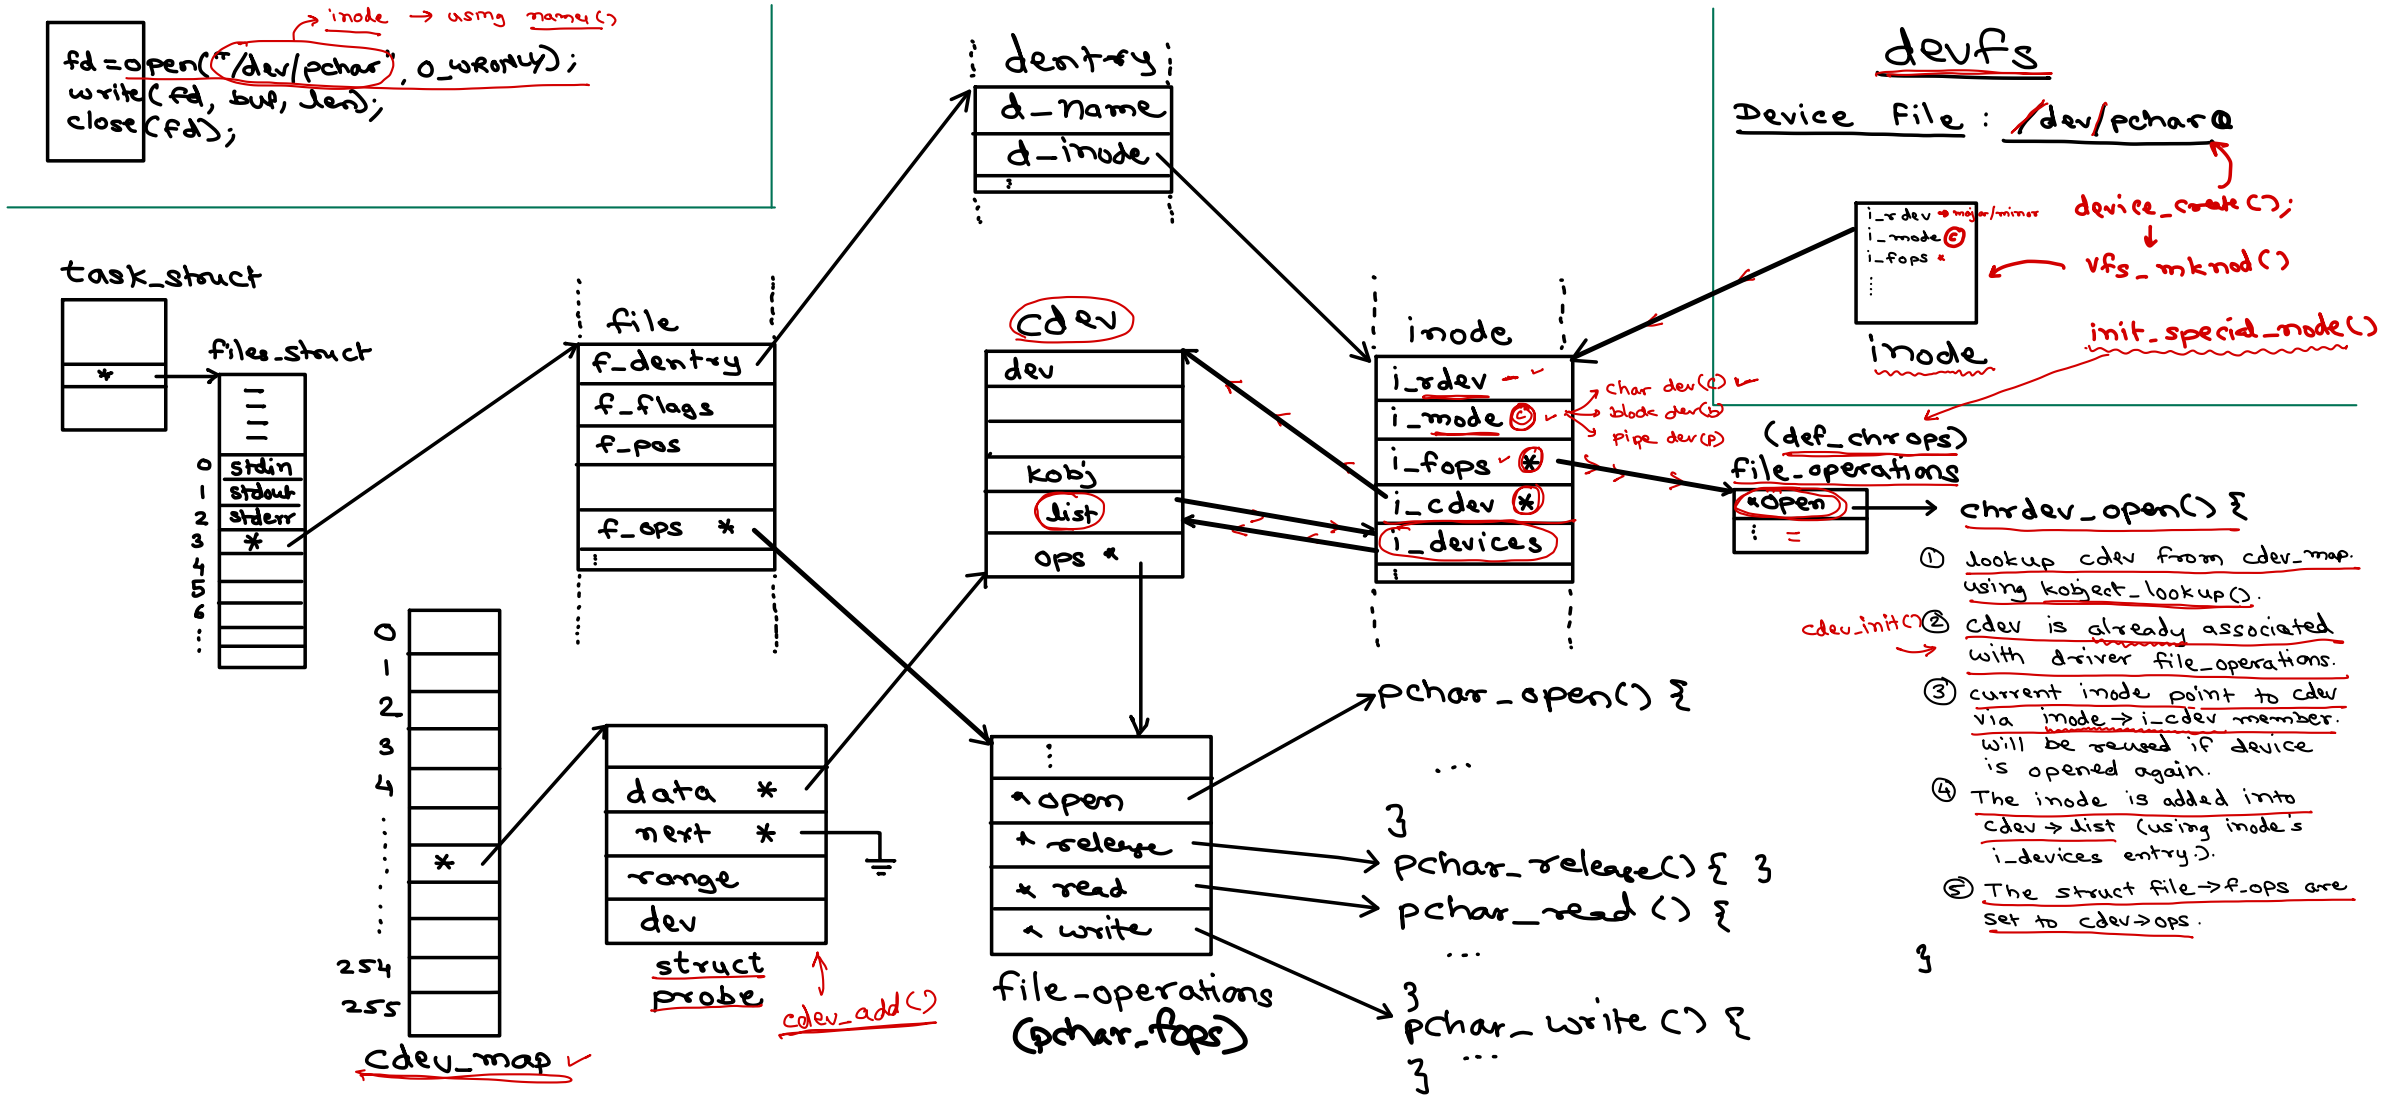

Waiting Queue

pchar device

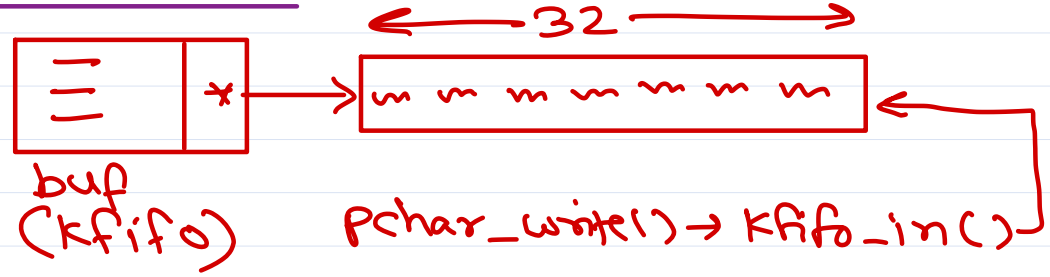

if kfifo is full, then block the writing process.

↳ wait_event_interruptible()

OR

wait_event()

↓

kfifo_is_full()

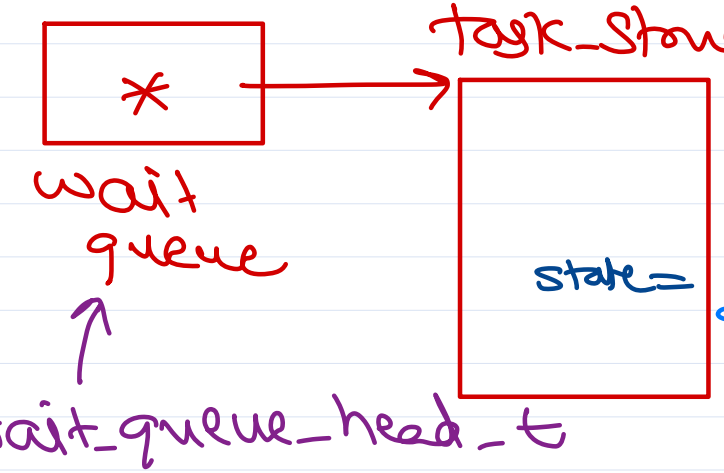

TASK_INTERRUPTIBLE → blocked until data is popped from fifo.
OR
Signal is received.

TASK_UNINTERRUPTIBLE → blocked until data is popped from fifo.

wake up blocked process, when some data is removed from fifo → kfifo_out() - read op.

OR
wake_up()
wake_up_interruptible()

Kernel stack - to create FARs of kernel functions

cmd> sudo insmod pchar.ko

insmod 2
Process

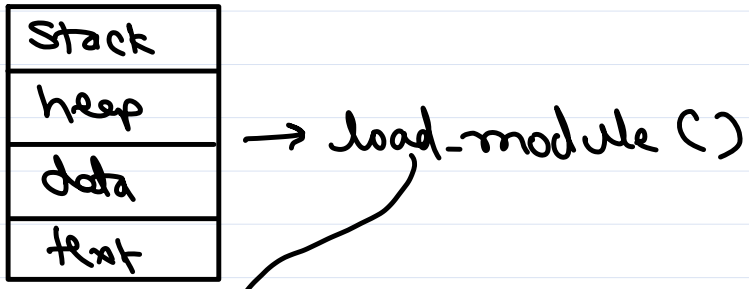

→ load_module()

sys_load_module()

{

① check privilege

② create struct module
state = COMING.

③ load module in kernel space.

④ call its init_module()

⑤ module state = LIVE;

}

task_struct (insmod)

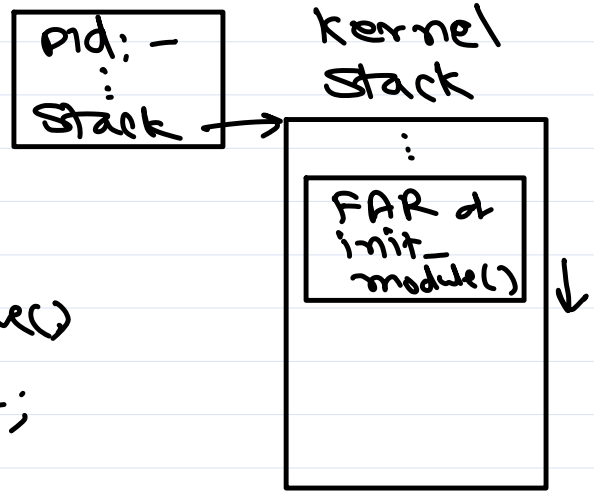

cmd> sudo cat /dev/pchar0

cat 2
Process

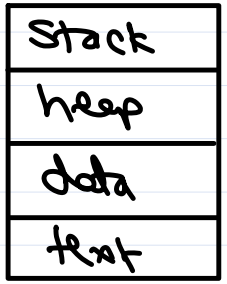

→ read()

sys_read() {

① get OFT entry
i.e. struct file

② call
file→f_ops→read()

}

pchar_read();

=

=

task_struct (cat)

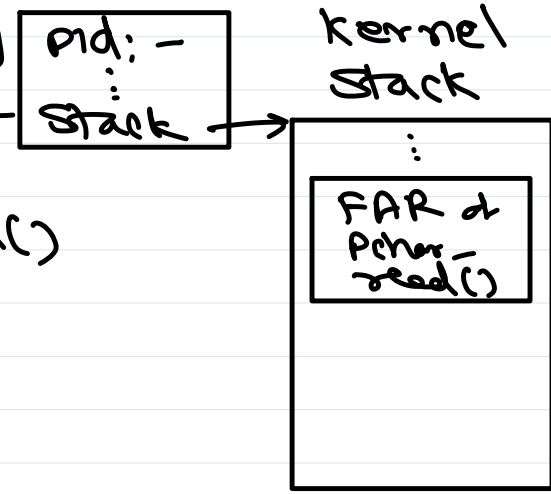

Get current process

* Each process (task) have a kernel stack.
  - on x86-32 arch, kernel stack size = 8KB.
  - the kernel stack is 8KB aligned i.e. its addr is multiple of 8KB (in kernel space).
  - task_struct has pointer to kernel stack.

* thread_info obj placed at the bottom of kernel stack (on x86).
  - it contains info about current thread & pointer to current task_struct.

* current_thread_info() returns addr of current task's thread_info object.
  → sp & -8192

* current macro returns current task_struct.
  #define current get_current()

  returns current process's task_struct address. ← current_thread_info()->task

task_struct and thread_info and kernel stack

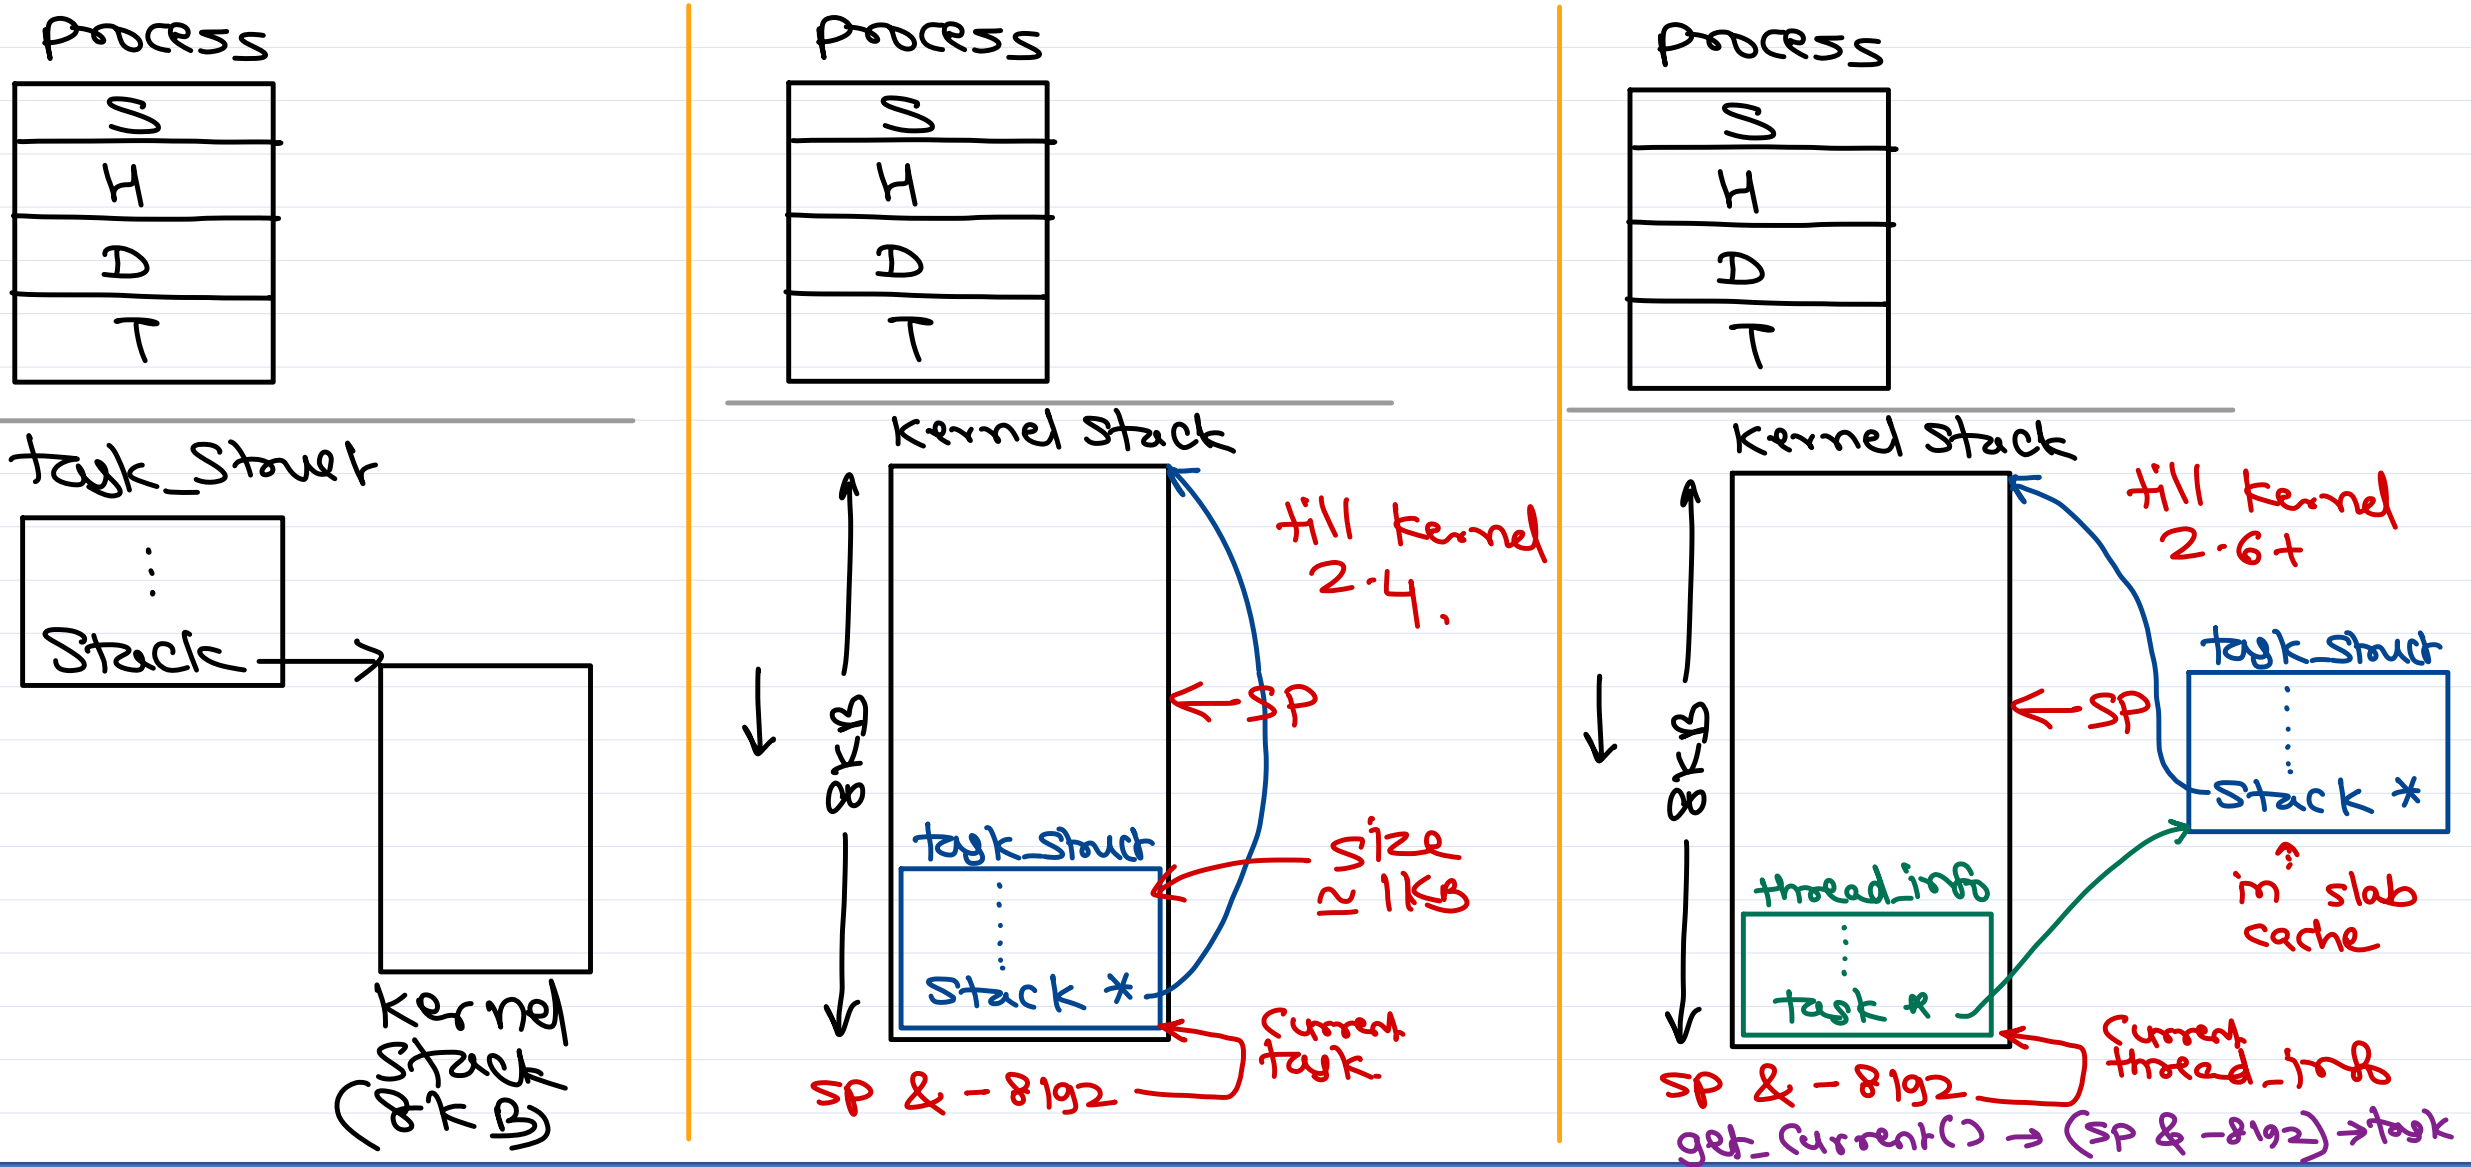

Synchronization in Kernel space

* multiple user processes may access kernel code (in drivers / syscalls /...) simultaneously and may cause race condition.
* kernel provided a few sync objects:
  1. semaphore.
  2. mutex (since 3.18+)
  3. spinlock

Semaphore in kernel space

<linux/semaphore.h>

semaphore

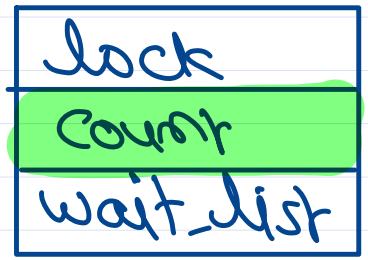

struct semaphore s;

init sem with given cnt
sema_init(&s, init cnt);

down(&s); → decr cnt & if -ve block process in uninterruptible sleep

down_interruptible(&s);
→ decr cnt & if -ve block process in interruptible sleep

up(&s); → incr count & wake up blocked process.

mutex in kernel space

3.18+ <linux/mutex.h>

mutex

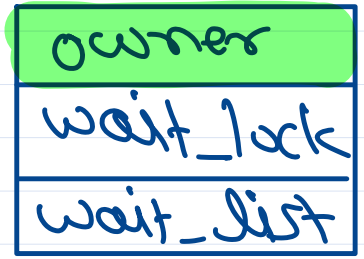

struct mutex m;

mutex_init(&m);

mutex_lock(&m);

mutex_lock_interruptible(&m);

mutex_unlock(&m);

mutex_destroy(&m);

Semaphore Internals

Solution 1: disable intrs

process 1

```
sem_p() {
    disable intrs;
    count--;
    ==
    enable intrs;
}
```

process 2

```
sem_v() {
    disable intrs;
    count++;
    ==
    enable intrs;
}
```

```
intr handler()
{
    ==
    scheduler();
    dispatcher();
}
```

Limitation: In multiprocessor env disabling intrs will disable them on current CPU only. Other process on other CPU can still cause the race condition.

Solution 2: spin locks

lock

process 1

```
sem_p() {
    spin_lock(&l);
    count--;
    ==
    spin_unlock(&l);
}
```

process 2

```
sem_v() {
    spin_lock(&l);
    count++;
    ==
    spin_unlock();
}
```

```
#include <linux/spinlock.h>
```

declare spinlock obj:

```
spinlock_t l;
```

init the spinlock

```
spin_lock_init(&l);
```

lock the spinlock:

```
spin_lock(&l);
```

unlock the spinlock:

```
spin_unlock(&l);
```

Spinlock

- hw based sync mechanism.
- it uses special instructions of exclusive access/bus holding.
  - ① ARM7 → SWP instruction.
    - Reference: Sloss book.
  - ② ARM cortex A/M →
    - LDREX, STREX, CLREX
    - Reference: Yiu book.
- these special instruction - test_and_set() kind.
  - testing (read) var value and setting var value is done in exclusive access to bus.
  - only one CPU can access bus at a time.

spinlock pseudo code:

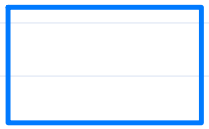

① spinlock_init():

```
lock = 0;
```

② spin_lock():

```
while(lock == 1)
    ;
lock = 1;
```

③ spin_unlock():

```
lock = 0;
```

ARM instructions:

LDREX: Load var value in CPU regr and mark that addr for exclusive access.

STREX: Store CPU regr value in given var. STREX will be successful only if it is done after LDREX for that addr.

CLREX: clear exclusive access for given addr (marked by LDREX).

addr of lock var in RAM


```
spinlock(lock) {
    do {
        while(__LDREX(lock) != 0)
            ;
        status = __STREX(1, lock);
    } while (status != 0);
    __DMB();
}
```

Spinlock → in Linux kernel.

```
#include <linux/spinlock.h>
spinlock_t lock;

init:
    spin_lock_init(&lock);
        ↳ lock=0;
lock:
    spin_lock(&lock);
        ↳ while(lock!=0)
            ;
          lock=1;
unlock:
    spin_unlock(&lock);
        ↳ lock=0;
```

Kernel debugging techniques

① debugging by printing

printk() → log levels @ 7
EMERG...DEBUG

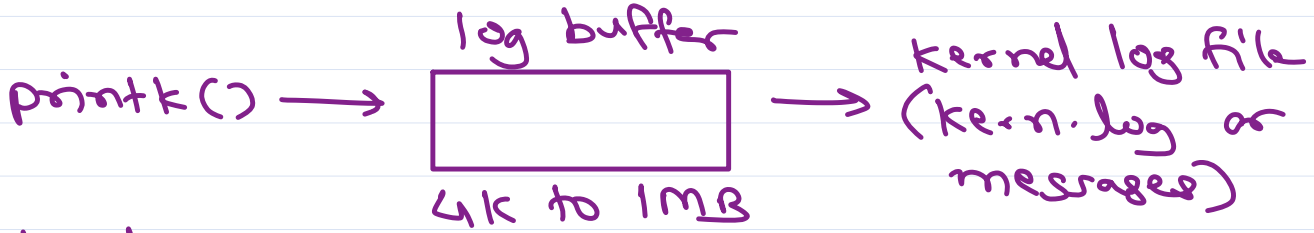

kernel log daemon → syslogd or klogd.

② debugging by querying

① ioctl() operation ② procfs entry.

③ debugging by watching

strace command → shows which sys calls are called

④ System faults/hangs

① Kernel OOPs message → register values + Stack trace.

② Kernel panic - Crash - Scheduler stopped

⑤ Kernel debuggers

① gdb ② kdb

Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>