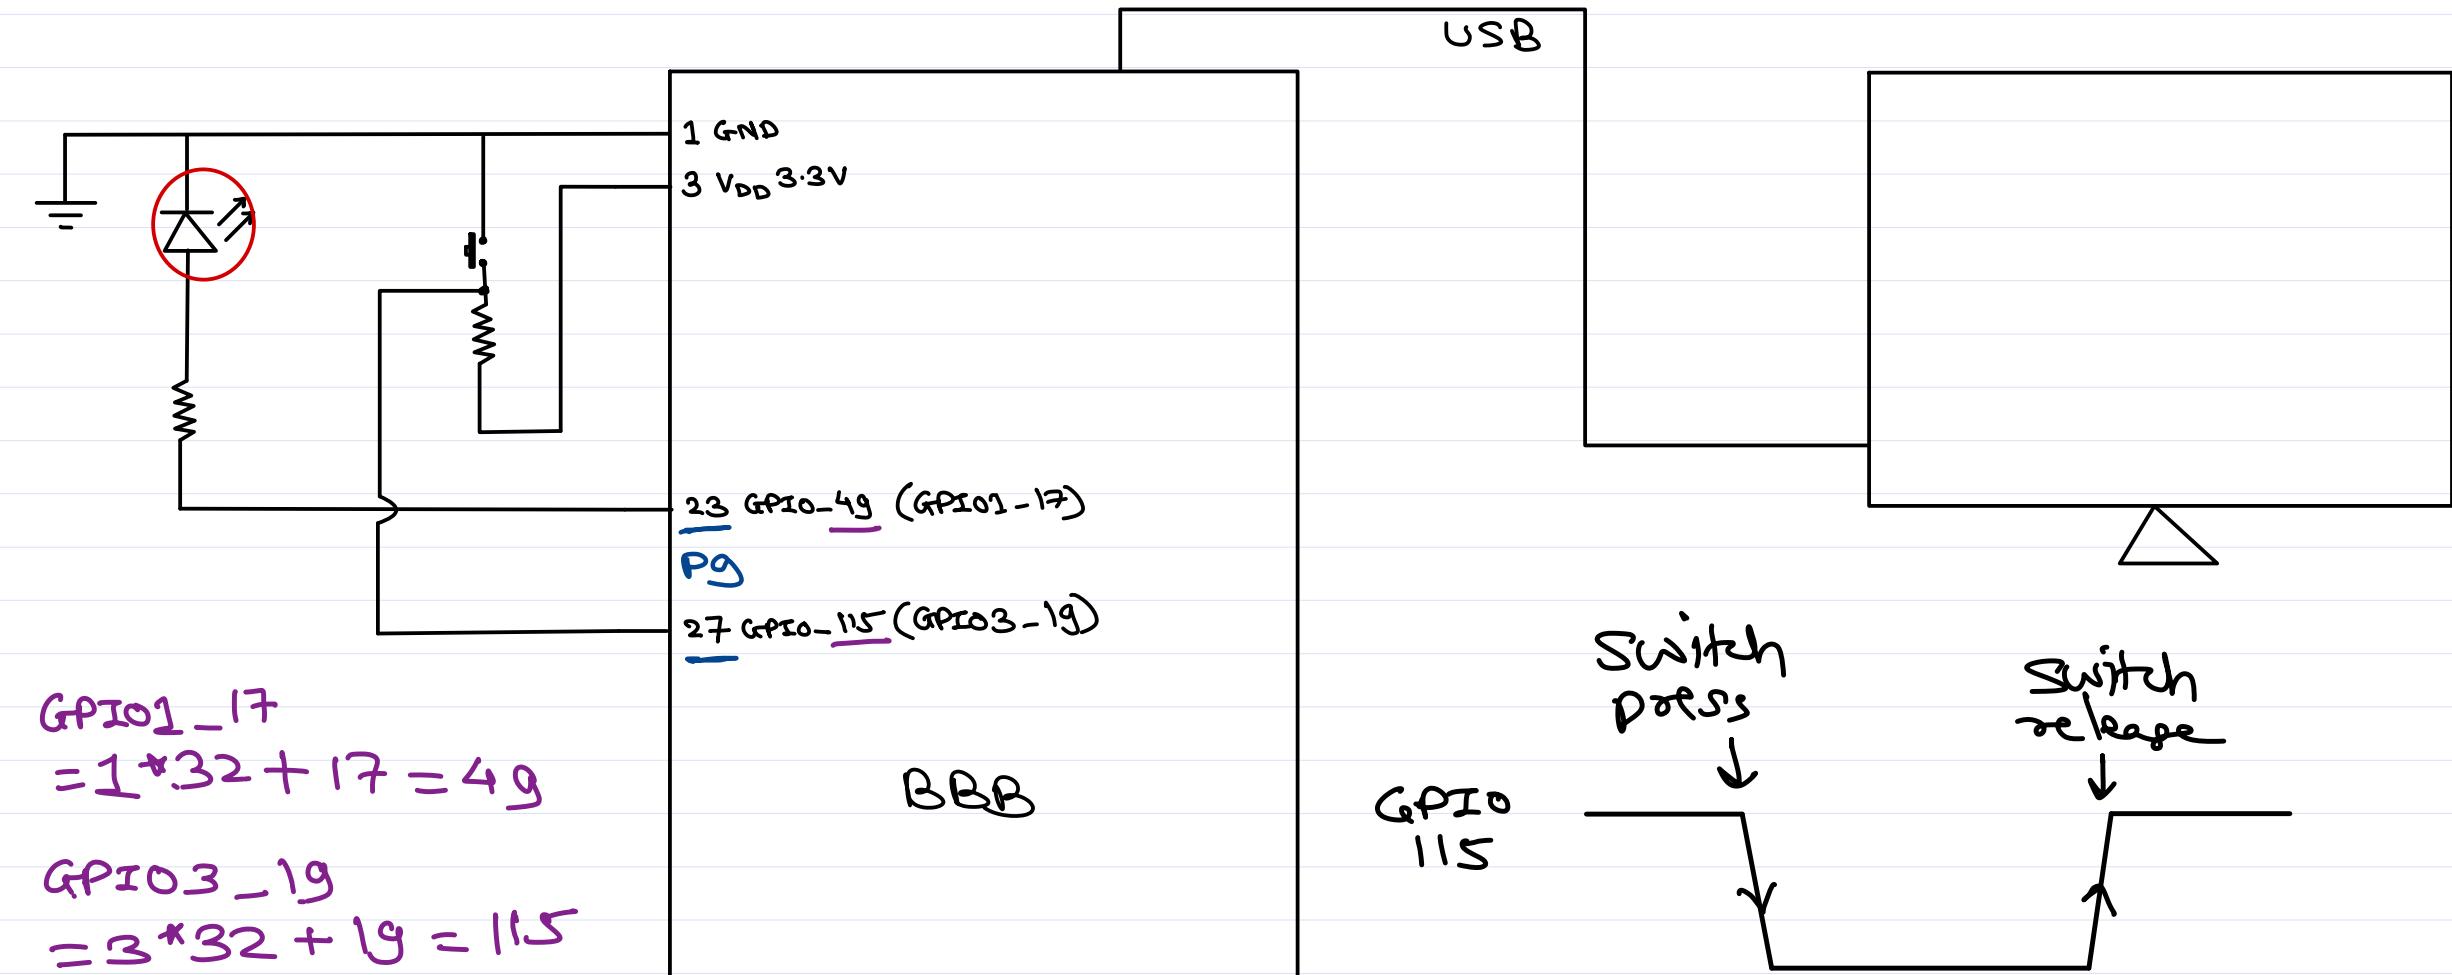


# Linux Character Device Driver

*Sunbeam Infotech*



# BBB Led & Switch

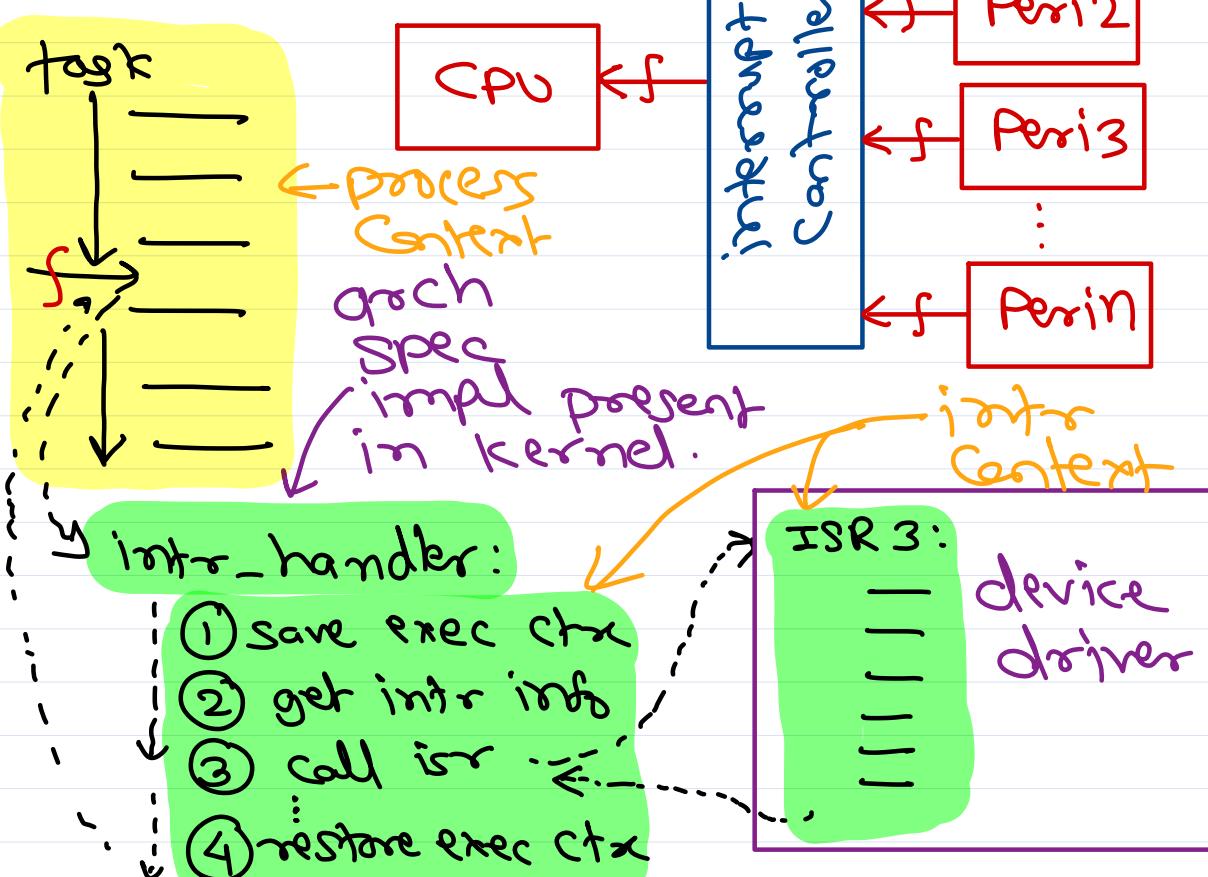


# Interrupt Handling in Linux

x86 → PIC - 8259

arm7 → VIC

arm cortex-m → NVIC



## ISR impl:

`irqreturn_t isr_fn(int irq, void *param);`

↳ IRQ\_HANDLED → intr handled by this ISR.  
↳ IRQ\_NONE → intr not handled by this ISR.  
↳ kernel calls next isr on same line

## Register ISR:

`request_irq(irq, isr_fn, flags, name, param);`

`cat /proc/interrupts`

- ① IRQF\_DISABLED
- ② IRQF\_TIMER
- ③ IRQF\_SAMPLE\_RANDOM ↳ true random numbers  
↳ /dev/random
- ④ IRQF\_SHARED ↳ random entropy pool  
↳ Same irq line is shared for multiple hw devices.

## Unregister ISR:

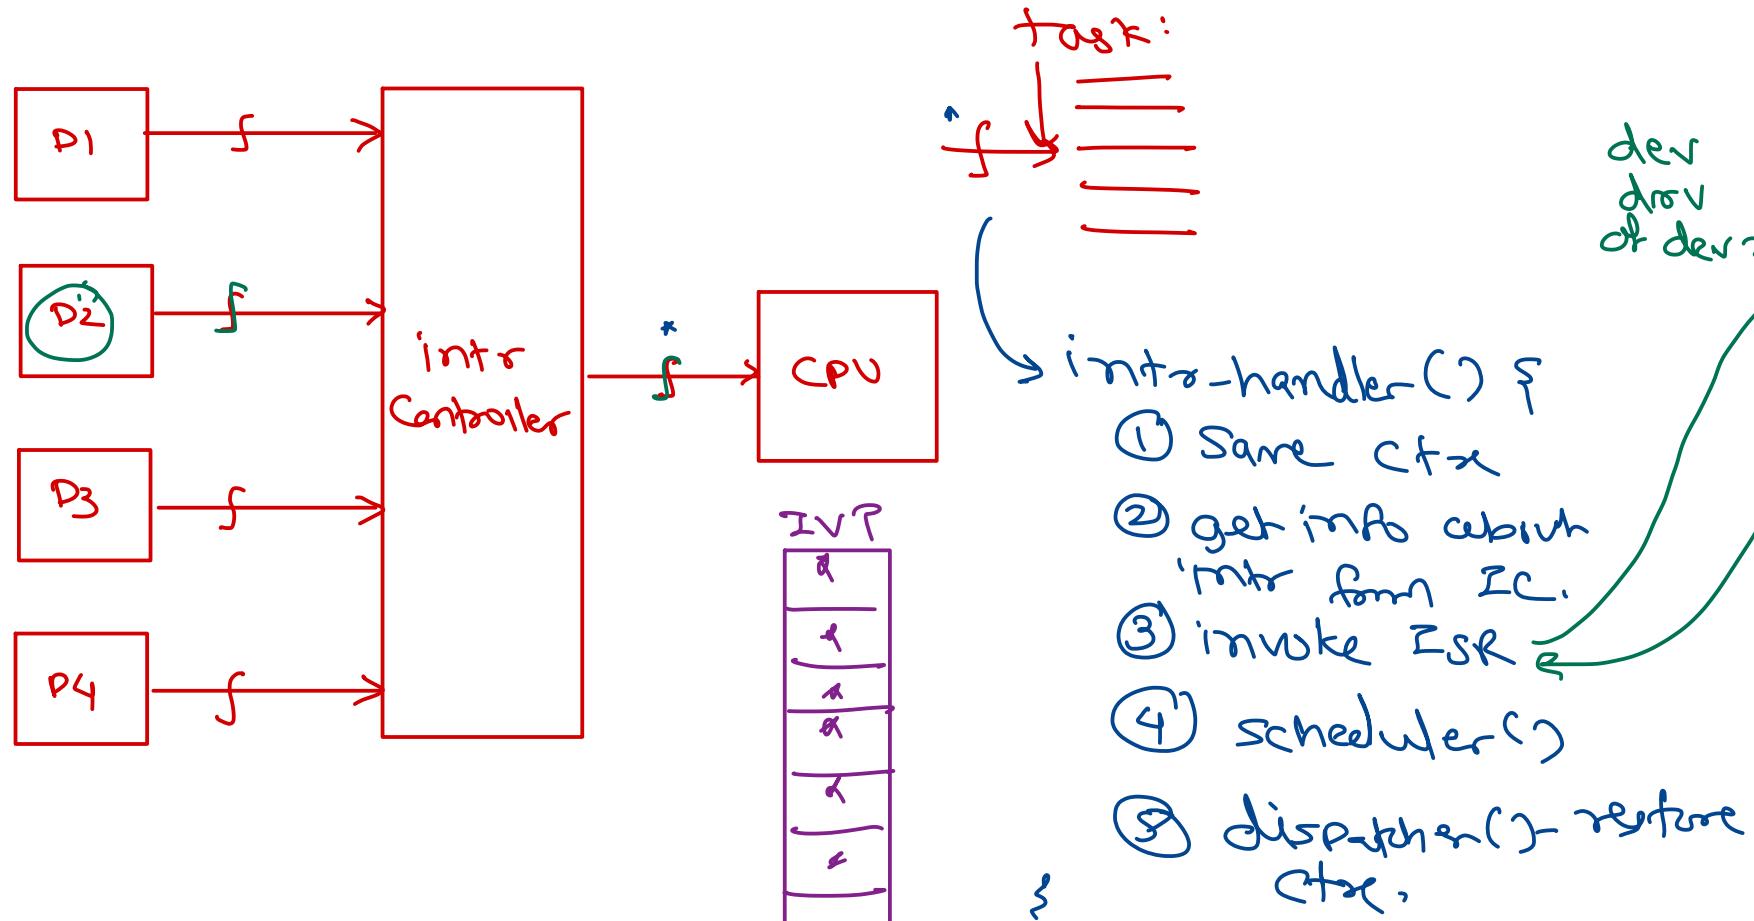
`free_irq(irq, param);`

e.g. com1, com2, com3, com4 ← Serial ports

4 3 4 3

# Interrupt

- Interrupts are special signals sent from device to CPU.
- Interrupt handling is architecture specific.



# Interrupt handling

- Interrupt is sent from the device to the PIC.
- PIC inform CPU about interrupt through interrupt line.
- CPU pause current task execution and execute interrupt handler.
- Interrupt handler does following
  - Save current task context on stack.
  - Get interrupt details from PIC.
  - Call ISR to handle the interrupt.
  - Invoke scheduler.
  - Restore the task context.
- In Linux there are two execution context.
  - Process context
    - User space process or kernel thread context. May block.
  - Interrupt context
    - Interrupt handler and ISR execution context.
    - Atomic context: cannot block.

uses user stack (n mb)  
or kernel stack (8 KB)

uses interrupt stack (<1 KB).  
per CPU



# Interrupt handling in Linux

- Since interrupt context cannot block, handler/ISR should return immediately.
- Heavy processing and/or blocking tasks should be deferred.
- Linux divides interrupt handling into two parts
  - Top half → ISR → Int Context
    - Run immediately when interrupt arrives.
    - Do time critical and non-blocking task like interrupt acknowledgement.
    - Cannot be pre-empted by another interrupt from same device.
  - Bottom half → Soft IRQ, Tasklet, Work Queue
    - Variety of bottom half implementations in Linux kernel.
    - Execute later – in interrupt context or process context.
    - Do heavy processing and/or blocking tasks.
    - Can be pre-empted by interrupt (top-half).
- Interrupt handling must be done in corresponding device driver.
  - Driver should implement top-half and/or bottom-half as per requirement.
  - Linux kernel ensure uniform programming model irrespective of architecture.



# Implementing top half

- Two step process
  - Implement ISR.
  - Register ISR.
- ISR registration
  - `#include <linux/interrupt.h>`
  - `int request_irq(unsigned int irq, irq_handler_t handler, unsigned long flags, const char *name, void *dev);`
    - irq: interrupt number
    - handler: typedef irqreturn\_t (\*irq\_handler\_t)(int, void \*);
    - flags:
      - IRQF\_DISABLED
      - IRQF\_SAMPLE\_RANDOM
      - IRQF\_TIMER
      - IRQF\_SHARED
    - name: device name /proc/irq and /proc/interrupts
    - dev: extra information to be passed to handler.
    - Returns 0 on success or -EBUSY if interrupt line is already in use.
  - request\_irq() may block and should not be called from interrupt context. Typically called when opening the device for processing or module initialization.

irqreturn\_t my\_isr(int irq, void \*param) {  
 =  
 3  
} =  
request\_irq()

multiple device instances - private obj  
to keep each device info  
e.g. struct serial\_info{  
 int irq;  
 int io\_addr;  
 mutex m;  
 ...  
};  
struct private\_struct  
devices[4];

com1  
3;  
request\_irq(4, my\_isr, IRQF\_SHARED, "com1", &devices[0]);

com3

request\_irq(4, my\_isr, IRQF\_SHARED, "com3", &devices[2]);

0	irq=4 io=0x3F8 : com1
1	irq=3 io=0x2F8 : com2
2	irq=4 io=0x3E8 : com3
3	irq=3 io=0x2E8 : com4



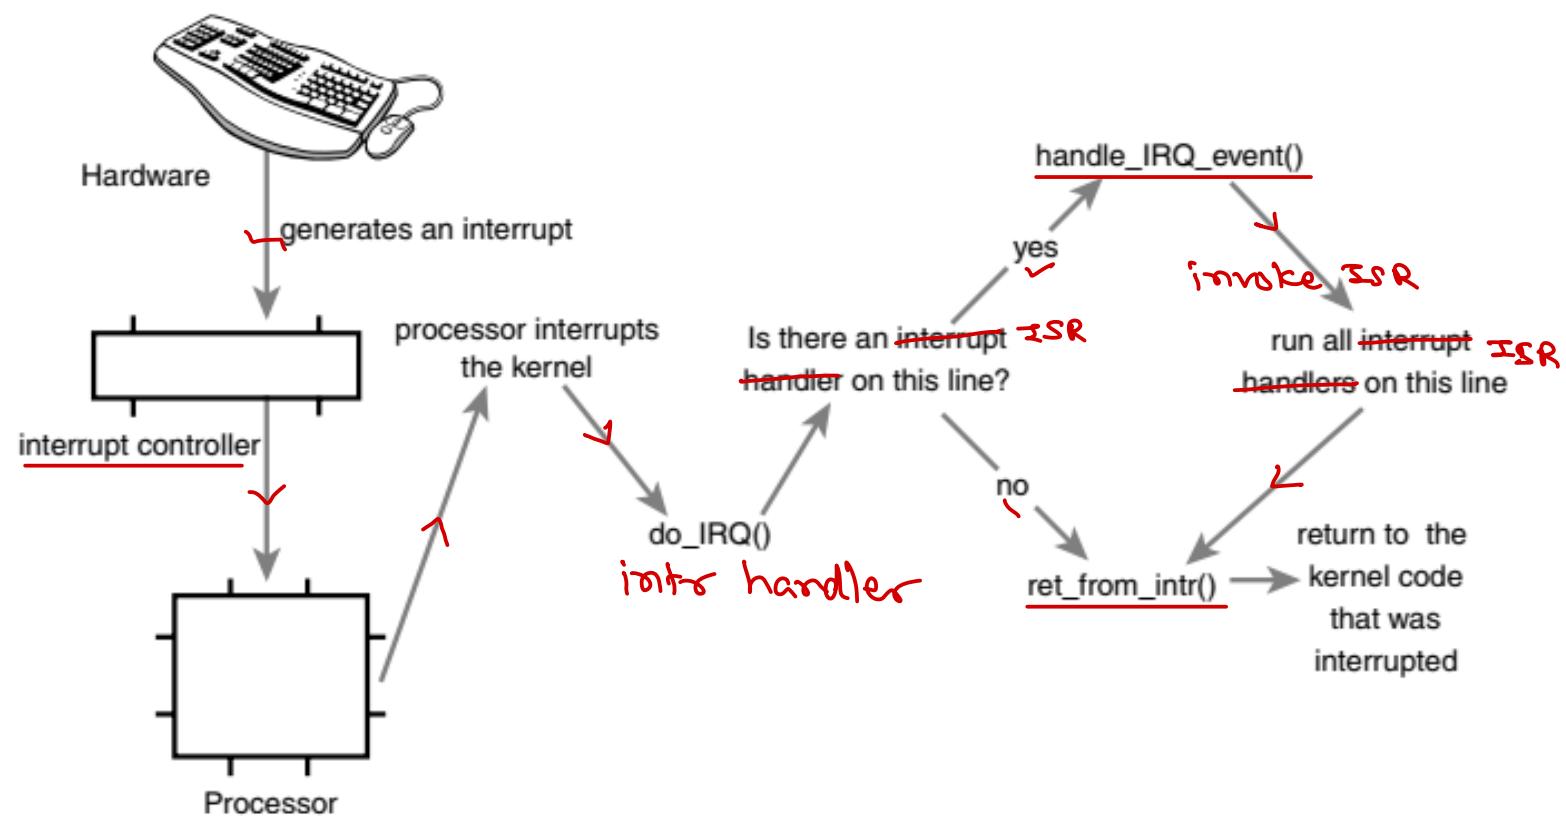
# Implementing top half

- ISR un-registration
  - Interrupt line must be released while unloading module or closing device.
  - void free\_irq(unsigned int irq, void \*dev);
- Implementing ISR
  - irqreturn\_t my\_intr\_handler(int irq, void \*dev);
    - irq: interrupt number
    - dev: extra param passed while request\_irq()
    - returns IRQ\_HANDLED or IRQ\_NONE.
  - Should contain time-critical tasks and interrupt acknowledgement.
  - Also trigger bottom-half if required.
  - Should not sleep/block.
- Linux interrupt handlers are not re-entrant. Current interrupt line is disabled while execution of ISR.
- Shared interrupt handlers
  - Must pass unique dev param – typically device private struct.
  - ISR must check if interrupt is raised from the corresponding device before handling it.
  - Kernel execute all ISR registered on same interrupt line.



# Interrupt handling

- Interrupt context
  - Atomic context.
  - One page kernel stack per processor.
- Interrupt execution



# Interrupt control

- ✗ • local\_irq\_disable(): Disables local interrupt delivery
- ✗ • local\_irq\_enable(): Enables local interrupt delivery
- ✓ • local\_irq\_save(): Saves the current state of local interrupt delivery and then disables it
- ✓ • local\_irq\_restore(): Restores local interrupt delivery to the given state
- disable\_irq(): Disables the given interrupt line and ensures no handler on the line is executing
- enable\_irq(): Enables the given interrupt line
- irqs\_disabled(): Returns nonzero if local interrupt delivery is disabled; otherwise returns zero
- in\_interrupt(): Returns nonzero if in interrupt context and zero if in process context
- in\_irq(): Returns nonzero if currently executing an interrupt handler and zero otherwise

ISR



# Bottom half

- If interrupt handling need to do heavy processing or blocking task, then driver must implement it in bottom half.
- General guideline for top and bottom half work division:
  - If the work is time sensitive, perform it in the interrupt handler. ISR
  - If the work is related to the hardware, perform it in the interrupt handler. ISR
  - If the work needs to ensure that another interrupt doesn't interrupt it, perform it in the interrupt handler. 2 ISR
  - For everything else, consider performing the work in the bottom half.
- Bottom half are executed after top half.
  - Immediately after top half in interrupt context.
  - In some specialized process context, when no other another high priority task is running.
- Types of bottom halves
  - BH - Removed in kernel 2.5
  - Task queue - Removed in kernel 2.5
  - Softirq - Added in kernel 2.3
  - Tasklet - Added in kernel 2.3
  - Work queue - Added in kernel 2.3



# Softirq

- Softirqs are statically allocated at compile time.

```
struct softirq_action {  
    void (*action)(struct softirq_action *);  
};  
static struct softirq_action softirq_vec[NR_SOFTIRQS];
```

→ 32

- Softirqs are implemented for specialized sub-systems.
- Kernel 2.6.34 have 9 Softirqs implemented.

• HI_SOFTIRQ	0	High-priority tasklets
• TIMER_SOFTIRQ	1	Timers
• NET_TX_SOFTIRQ	2	Send network packets
• NET_RX_SOFTIRQ	3	Receive network packets
• BLOCK_SOFTIRQ	4	Block devices
• TASKLET_SOFTIRQ	5	Normal priority tasklets
• SCHED_SOFTIRQ	6	Scheduler
• HRTIMER_SOFTIRQ	7	High-resolution timers
• RCU_SOFTIRQ	8	RCU locking



# Softirqs

- Softirqs must be triggered for the execution. This is called as “raising Softirq”. Mostly done from ISR.
- Pending Softirq are checked and executed in one of the following place – do\_softirq().
  - In the return from hardware interrupt code path *← interrupt context*
  - In the ksoftirqd kernel thread (per processor) *← process context*
  - In any code that explicitly checks for and executes pending softirqs, such as the networking subsystem
- Using Softirq
  - Currently only network and block subsystem is using Softirq directly.
  - It is not advised to use Softirqs directly.
  - Softirq can be registered using open\_softirq() and can be triggered using raise\_softirq().



# Tasklets

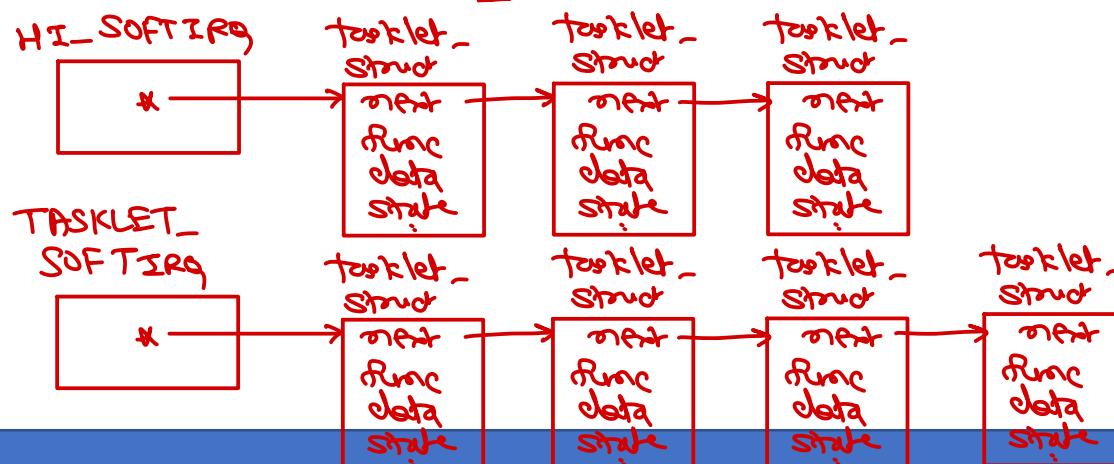
①

5

- Implemented on top of Softirqs i.e. HI\_SOFTIRQ and TASKLET\_SOFTIRQ.
- Tasklets are dynamic components and much easier to use.

```
struct tasklet_struct {  
    struct tasklet_struct *next; /* next tasklet in the list */  
    unsigned long state; /* state of the tasklet */  
    atomic_t count; /* reference counter */  
    void (*func)(unsigned long); /* tasklet handler function */  
    unsigned long data; /* argument to the tasklet function */  
};
```

- Tasklet state can be: 0, TASKLET\_STATE\_SCHED or TASKLET\_STATE\_RUN.



# Tasklets

- Declare Tasklet statically.
  - `DECLARE_TASKLET(my_tasklet, my_tasklet_handler, dev);`
- Declare & initialize Tasklet dynamically.
  - struct tasklet\_struct my\_tasklet;
  - tasklet\_init(t, tasklet\_handler, dev); *arg to tasklet func.*
- Tasklet handler implementation
  - `void tasklet_handler(unsigned long data) { ... }`
  - Like softirq, tasklet cannot sleep/block.
  - While executing tasklet, interrupts are enabled.
  - Tasklet are not re-entrant or execute concurrently.
- Trigger Tasklet
  - tasklet\_schedule(&my\_tasklet); *tasklet\_hi\_schedule(&my\_tasklet);*
  - Change tasklet state to TASKLET\_STATE\_SCHED.
- Tasklet can be enabled/disabled explicitly.
  - tasklet\_enable(&my\_tasklet);
  - tasklet\_disable(&my\_tasklet);



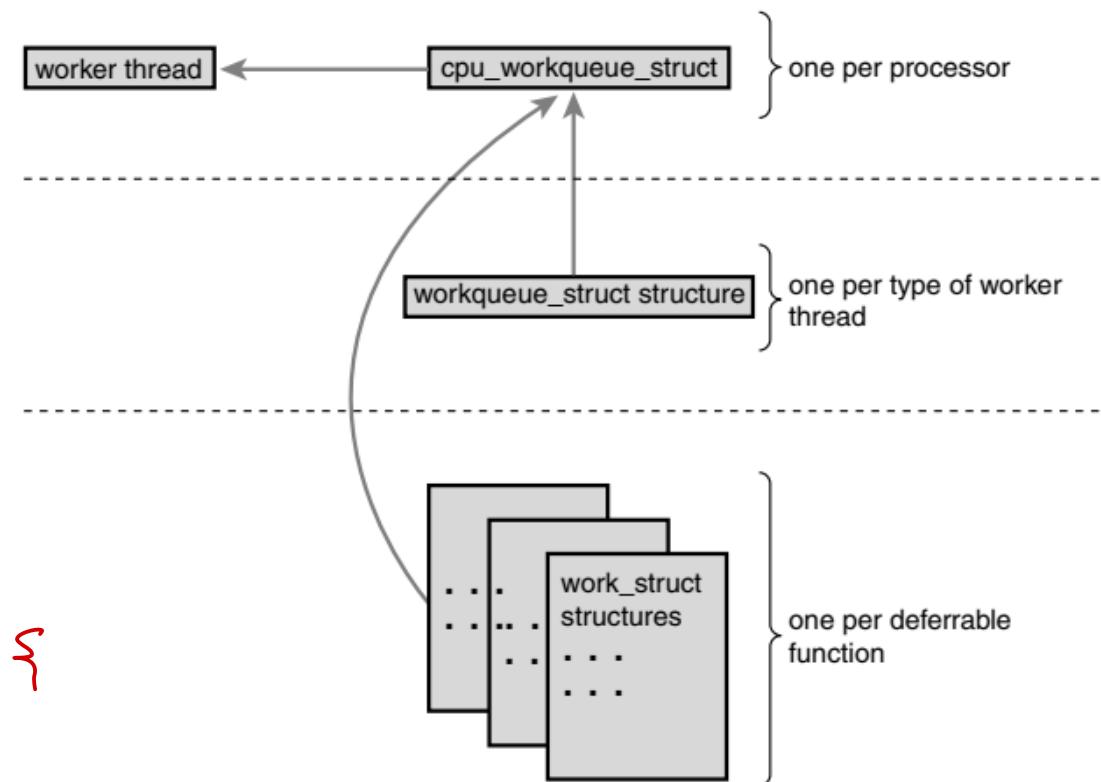
# Work queue

- Work queues defer work into a kernel thread. *→ kworker*
- Always runs in process context – worker threads (per processor).
- Work queues are schedulable and can sleep/block.
- Usual alternative to work queues is kernel threads. However creating new kernel threads isn't advised.

```
struct cpu_workqueue_struct {  
    spinlock_t lock; /* lock protecting this structure */  
    struct list_head worklist; /* list of work */  
    wait_queue_head_t more_work;  
    struct work_struct *current_struct;  
    struct workqueue_struct *wq; /* associated */  
    task_t *thread; /* associated thread */  
};  
  
struct work_struct {  
    atomic_long_t data;  
    struct list_head entry;  
    work_func_t func;  
};
```

*wid\_func( data ) {*

*work queue* → shared (by kernel)  
→ dedicated (by driver).



# Work queue (*in shared work queue*)

- Creating work

- DECLARE\_WORK(name, void (\*func)(void \*), void \*data); // static
- INIT\_WORK(struct work\_struct \*work, void (\*func)(void \*), void \*data); // dynamic

- Work handler

- void work\_handler(void \*data) { ... } → may sleep

- Scheduling work

- schedule\_work(&work);
- schedule\_delayed\_work(&work, delay);

} in shared  
work queue

- Ensure work completion

- void flush\_scheduled\_work();

- Cancel scheduled work

- int cancel\_delayed\_work(struct work\_struct \*work);



# Control bottom half

- To enable/disable all bottom half (tasklet & work queue) processing
  - void local\_bh\_enable();
  - void local\_bh\_disable();
- Choosing bottom half
  - Softirq - **Static component**
    - Concurrent execution and need synchronization
    - Good for fast execution and high frequency use
    - Cannot sleep (*may execute in process ctx-kssoftirq or may execute in intr ctx-return from 'intr'*)
  - Tasklet - **dynamic Component**
    - Simplified programming
    - Not executed concurrently - *no sync needed*
    - Cannot sleep - *internally depend on softirq.*
  - Work queue - **dynamic Component**
    - Can sleep - *runs in process ctx-kworker.*
    - Higher overheads due to kernel thread & context switching

kworker





*Thank you!*

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>

