

# Embedded Operating Systems

## Agenda

- Directory
- Links
- File IO System Calls
  - open()
  - close()
  - read()
  - write()
  - lseek()
- Disk allocation mechanisms
- Free space management mechanisms

## Current Directory

- Each process has a current directory.
- The relative paths in the program are processed w.r.t. the current directory of that process.
- In Linux, struct task\_struct (PCB) has a member struct fs\_struct.

```
struct fs_struct {  
    struct dentry *pwd;  
    struct dentry *root;  
    // ...  
};
```

- By default current directory is inherited from the parent process.

## chdir(dirpath)

- Change the current directory of the current process to the given directory path.
- arg1: dir path
- returns: 0 on success.

## Directory permissions

- r (read): Can read directory entries.
  - ls command can list the directory.
  - readdir() can get directory entry.
- w (write): Can add new directory entry, modify existing directory entry, or delete directory entry.
  - new directory entry: when create new sub-directory (e.g. mkdir) or file (e.g. touch, ...).
  - modify existing directory entry: rename file or sub-directory.
  - delete directory entry: delete file or sub-directory (e.g. rm, rmdir, ...)
- x (execute): Can browse the directory.

- cd command (chdir() syscall) will work

## Links

- Links are shortcuts to access deeply located files quickly.
- There are two types of links in Linux/UNIX.
  1. Symbolic Link
  2. Hard Link

### Symbolic Link

- terminal> ln -s /path/of/target/file linkpath

```
cd ~
mkdir links
cd links
echo "file one" > one.txt
cat one.txt
ls -l -i
ln -s one.txt two.txt
echo "link created" >> two.txt
cat two.txt
cat one.txt
ln -s /home/sunbeam/links/one.txt three.txt
cat three.txt
ls -l -i
rm one.txt
ls -l -i
cat two.txt
rm two.txt three.txt
```

- Internally use symlink() syscall.
  - man symlink
  - int symlink(const char \*target\_path, const char \*link\_path);

### ~~SUNBEAM INFO TECH~~symlink() syscall

- A new link file is created (new inode and new data block is allocated), which contains info about the target file (absolute or relative path).
- Link count is not incremented.
- If target file is deleted, the link becomes useless.
- Can create symlinks for directories also.

### Hard Link

- terminal> ln targetfilepath

```
pwd
echo "file one." > one.txt
cat one.txt
ls -l -i
ln one.txt two.txt
echo "hard link created" >> two.txt
cat two.txt
cat one.txt
ls -l -i
ln one.txt three.txt
ls -l -i
rm one.txt
ls -l -i
cat two.txt
```

- Internally use link() syscall.
  - man link
  - int link(const char \*target\_path, const char \*link\_path);

### ~~link()~~ syscall

- A new directory entry is created, which has a new name and same inode number. No new file (inode and data blocks) is created.
- Link count in the inode of the file is incremented.
- If directory entry of target file is deleted (rm command), file can be still accessed by link directory entry.
- Cannot create hard link for directories, because it may lead to infinite recursion (while traversing directories recursively e.g. ls -R)

### ~~rm~~ command

- The rm command in Linux, internally calls unlink() system call.
  - int unlink(const char \*filepath);

### ~~unlink()~~ syscall

- It deletes directory entry of the file.
- It decrements link count in the inode by 1.
- If link count = 0, the inode is considered to be deleted/free (updated into super-block). It can be reused for any new file.
- When inode is marked free, data blocks are also made free, so that they can also be reused for some new file.

## File IO syscalls

### open() syscall

- fd = open("file-path", flags, mode);

- arg1: path of file to be opened
- arg2: flags - how you want to open the file
  - O\_RDONLY, O\_WRONLY, O\_RDWR -- read-write flags
  - O\_TRUNC -- delete the contents of file while opening
  - O\_APPEND -- write at the end of file
  - O\_CREAT -- create a new file (if not present) -- must give arg3
- arg3: mode - permissions for new file - octal number
- returns file descriptor on success, and -1 on failure.
  - file descriptor is int that uniquely identifies the file in that process.
  - fd is used in other file io syscalls e.g. close(), read(), write(), lseek().
- MCQ questions
  - fopen "w" mode is equivalent to O\_WRONLY | O\_TRUNC | O\_CREAT
  - fopen "r" mode is equivalent to ...
  - fopen "a" mode is equivalent to ...
  - fopen "w+" mode is equivalent to O\_RDWR | O\_TRUNC | O\_CREAT
  - fopen "r+" mode is equivalent to ...
  - fopen "a+" mode is equivalent to ...
- fd = open("/home/nilesh/abc.txt", O\_RDONLY);
  - step 1. Convert given file path into its inode number. This is called as path name translation and is done by a kernel function namei().
  - step 2. Load the inode of the file from the disk into inode table in memory. Inodes of all recently accessed files are kept in this table.
  - step 3. A file position is initialized to 0 and is stored in the open file table. It also stores flags in which file is opened and pointer to the in-memory inode. Information of all files opened in the system, is maintained in this table.
  - step 4. Each process is associated with an open file descriptor table. It keeps info of all files opened by that process. This entry stores pointer to the OFT entry.
  - step 5. Finally index to file desc table entry is returned, which is called as "file descriptor". All further read(), write(), lseek(), close() operations will be using this file desc.

## VFS Structures

### **STRUCT INODE**

- unsigned long i\_ino; // inode number
- loff\_t i\_size; // file size
- unsigned int i\_nlink; // number of hard links
- umode\_t i\_mode; // file mode (permissions)
- atomic\_t i\_count; // reference count
- struct list\_head i\_list; // inode cache
- Device driver related
  - struct list\_head i\_devices;
  - dev\_t i\_rdev;
  - union { struct pipe\_inode\_info \*i\_pipe; struct block\_device \*i\_bdev; struct cdev \*i\_cdev; };
  - struct file\_operations \*i\_fop;

### **STRUCT FILE**

- unsigned int f\_flags; // open() arg2
- loff\_t f\_pos; // current file position
- struct path f\_path; // pointer to dentry
- #define f\_dentry f\_path.dentry
- struct list\_head fu\_list; // open file table
- atomic\_t f\_count; // reference count
- Device driver related
  - struct file\_operations \*f\_op;

### **struct dentry**

- struct qstr d\_name; // name of file/sub-directory
- struct inode \*d\_inode; // pointer to the inode
- struct list\_head d\_lru; // dentry cache
- atomic\_t d\_count; // reference count

### **struct fs\_struct**

- struct dentry \* root; // stores "root directory" of the process --> used for absolute path
- struct dentry \* pwd; // stores "current directory" of the process --> used for relative path
- int umask; // user file mode mask -- while creating new file this mask is used.

### **struct files\_struct**

- struct file \* fd\_array[NR\_OPEN\_DEFAULT];

### **struct task\_struct**

- struct fs\_struct \*fs; // current & root directory
- struct files\_struct \*files; // open file desc tables

### **Reference counting**

- Used to manage life-time of any object (in complex systems e.g. Linux kernel, ...).
- Object has a member called as "reference count". The count is incremented everytime new pointer points to the same object and decremented everytime the pointer to the object is no more used/required.
- At any moment, reference count is number of pointers referring to the object.
- When reference count become zero, it means no pointer is referring to the object and the object can be deleted safely.

### **File IO syscalls**

#### **close() syscall**

- Decrement ref count in open file table entry (struct file).
- If ref count drops to zero, OFT entry is deleted (from OFT).

#### **read() syscall**

- count = read(fd, buf, length); -- syscall api
  - sys\_read(fd, buffer, length) -- syscall implementation
  - vfs\_read(file, buffer, length, inode) -- Virtual file system
    - Logical FS considers file as sequential set of bytes and set of blocks.
    - Example: block size = 4096 and file size = 20000 bytes, then number of blocks = 5 (0 to 4).
      - If current file position = 10000, then reading file block = 2
  - ext3\_read(file, inode, file\_block) -- File system manager
    - Refers inode and file disk block corresponding to the file block.
    - check buffer cache -- if disk block found.
      - if found, return it;
      - otherwise call disk driver to read that disk block from disk
  - disk\_read(disk\_device, disk\_block) -- device driver
    - Read appropriate sectors and made it available into buffer cache.
    - The current process is blocked/sleep while disk read operation is in progress.