




**Sunbeam Institute of Information Technology
Pune and Karad**

Module - Embedded Operating System

Trainer - Devendra Dhande
Email – devendra.dhande@sunbeaminfo.com




Virtual file system



Annotations for the VFS layers:

- `fopen()`, `fclose()`, `fread()`, `fwrite()` → System Call API
- `Open()`, `read()`, `write()`, `close()` → System Call layer
- `sgs_open()`, `sgs_read()`, `sgs_write()`, `sys_close()` → Virtual File System
- `mount table`: `Partition`: `File System`
VFS redirects FS request to respective file system manager
VFS defines many struct/object to deal with file systems
 - superblock, file, inode, dentry
 - superblock-ops, file-ops, inode-ops, dentry-ops
- `manages FS layouts on partitions`
`super block, inode list, data blocks`
`read/write`
- `Schedules h/w access requests (request queue)`
- `cache the data of storage device`

Linux: ext2/3/4, NTFS...
Windows: NTFS, CDFS, FAT



Disk blocks:

- manufacturer → sector = 512B
- Filesystems → FS Block = 1K, 2K, 4K, 8K.... 2M



```
struct super_block {
    struct list_head    s_list;          /* list of all superblocks */
    dev_t               s_dev;           /* identifier */
    unsigned long       s_blocksize;     /* block size in bytes */
    unsigned long       s_old_blocksize; /* old block size in bytes */
    unsigned char       s_blocksize_bits; /* block size in bits */
    unsigned char       s_dirt;          /* dirty flag */
    unsigned long long  s_maxbytes;     /* max file size */
    struct file_system_type s_type;     /* filesystem type */
    struct super_operations s_op;       /* superblock methods */
    struct dquot_operations *dq_op;     /* quota methods */
    struct quotactl_ops  *s_qcop;        /* quota control methods */
    struct export_operations *s_export_op; /* export methods */
    unsigned long        s_flags;         /* mount flags */
    unsigned long        s_magic;         /* filesystem's magic number */
    struct dentry        *s_root;         /* directory mount point */
    struct rw_semaphore  s_umount;       /* unmount semaphore */
    struct semaphore     s_lock;          /* superblock semaphore */
    int                 s_count;         /* superblock ref count */
    int                 s_syncing;        /* filesystem syncing flag */
    int                 s_need_sync_fs;  /* not-yet-synced flag */
    atomic_t             s_active;        /* active reference count */
    void                *s_security;     /* security module */
    struct list_head    s_dirty;         /* list of dirty inodes */
    struct list_head    s_io;            /* list of writebacks */
    struct hlist_head   s_anon;          /* anonymous dentries */
    struct list_head    s_files;         /* list of assigned files */
    struct block_device *s_bdev;         /* associated block device */
    struct list_head    s_instances;     /* instances of this fs */
    struct quota_info   s_dquot;         /* quota-specific options */
    char                s_id[32];        /* text name */
    void                *s_fs_info;       /* filesystem-specific info */
    struct semaphore    s_vfs_rename_sem; /* rename semaphore */
};
```

```
struct super_operations {
    struct inode *(*alloc_inode) (struct super_block *sb);
    void (*destroy_inode) (struct inode *);
    void (*read_inode) (struct inode *);
    void (*dirty_inode) (struct inode *);
    void (*write_inode) (struct inode *, int);
    void (*put_inode) (struct inode *);
    void (*drop_inode) (struct inode *);
    void (*delete_inode) (struct inode *);
    void (*put_super) (struct super_block *);
    void (*write_super) (struct super_block *);
    int (*sync_fs) (struct super_block *, int);
    void (*write_super_lockfs) (struct super_block *);
    void (*unlockfs) (struct super_block *);
    int (*statfs) (struct super_block *, struct statfs *);
    int (*remount_fs) (struct super_block *, int *, char *);
    void (*clear_inode) (struct inode *);
    void (*umount_begin) (struct super_block *);
    int (*show_options) (struct seq_file *, struct vfsmount *);
};
```



VFS structures

```
struct inode {
    struct hlist_node    i_hash;
    struct list_head     i_list; ✓
    struct list_head     i_dentry; ✓
    unsigned long         i_ino; ✓
    atomic_t              i_count; ✓
    umode_t                i_mode; ✓
    unsigned int           i_nlink; ✓
    uid_t                  i_uid; ✓
    gid_t                  i_gid; ✓
    kdev_t                 i_rdev;
    loff_t                  i_size; ✓
    struct timespec        i_atime; ?
    struct timespec        i_mtime;
    struct timespec        i_ctime;
    unsigned int           i_blkbits;
    unsigned long           i_blksize;
    unsigned long           i_version;
    unsigned long           i_blocks;
    unsigned short          i_bytes;
    spinlock_t              i_lock;
    struct rw_semaphore     i_alloc_sem;
    struct semaphore        i_sem;
    struct inode_operations *i_op;
    struct file_operations *i_fop;
    struct super_block      *i_sb;
    struct file_lock        *i_flock;
    struct address_space     *i_mapping;
    struct address_space     i_data;
    struct dquot            *i_dquot[MAXQUOTAS];
    struct list_head         i_devices;
    struct pipe_inode_info   *i_pipe;
    struct block_device      *i_bdev;
    unsigned long            i_dnotify_mask;
    struct dnotify_struct    *i_dnotify;
    unsigned long            i_state;
    unsigned long            dirtied_when;
    unsigned int              i_flags;
    struct socket            i_sock;
    unsigned int              i_writecount;
    struct security_struct   *i_security;
    unsigned long            i_generation;
    void                     *generic_ip; /* filesystem-specific info */
    union {
        void
    } u;
};
```

```
struct inode_operations {
    int (*create) (struct inode *, struct dentry *, int);
    struct dentry * (*lookup) (struct inode *, struct dentry *);
    int (*link) (struct dentry *, struct inode *, struct dentry *);
    int (*unlink) (struct inode *, struct dentry *);
    int (*symlink) (struct inode *, struct dentry *, const char *);
    int (*mkdir) (struct inode *, struct dentry *, int);
    int (*rmdir) (struct inode *, struct dentry *);
    int (*mknod) (struct inode *, struct dentry *, int, dev_t);
    int (*rename) (struct inode *, struct dentry *,
                  struct inode *, struct dentry *);
    int (*readlink) (struct dentry *, char *, int);
    int (*follow_link) (struct dentry *, struct nameidata *);
    int (*put_link) (struct dentry *, struct nameidata *);
    void (*truncate) (struct inode *);
    int (*permission) (struct inode *, int);
    int (*setattr) (struct dentry *, struct iattr *);
    int (*getattr) (struct vfsmount *, struct dentry *, struct kstat *);
    int (*setxattr) (struct dentry *, const char *,
                    const void *, size_t, int);
    ssize_t (*getxattr) (struct dentry *, const char *, void *, size_t);
    ssize_t (*listxattr) (struct dentry *, char *, size_t);
    int (*removexattr) (struct dentry *, const char *);
};
```





VFS structures

```
struct dentry {  
    atomic_t          d_count;      /* usage count */  
    unsigned long     d_vfs_flags;  /* dentry cache flags */  
    spinlock_t        d_lock;       /* per-dentry lock */  
    struct inode      *d_inode;     /* associated inode */  
    struct list_head  d_lru;        /* unused list */  
    struct list_head  d_child;     /* list of dentries within */  
    struct list_head  d_subdirs;   /* subdirectories */  
    struct list_head  d_alias;     /* list of alias inodes */  
    unsigned long     d_time;       /* revalidate time */  
    struct dentry_operations *d_op; /* dentry operations table */  
    struct super_block *d_sb;       /* superblock of file */  
    unsigned int       d_flags;      /* dentry flags */  
    int               d_mounted;    /* is this a mount point */  
    void *            d_fsdta;     /* filesystem-specific data */  
    struct rcu_head    d_rcu;       /* RCU locking */  
    struct dcookie_struct *d_cookie; /* cookie */  
    struct dentry     *d_parent;    /* dentry object of parent */  
    struct qstr        d_name;      /* dentry name */  
    struct hlist_node  d_hash;      /* list of hash table entries */  
    struct hlist_head  *d_bucket;   /* hash bucket */  
    unsigned char      d_iname[DNAME_INLINE_LEN_MIN]; /* short name */  
};
```

```
struct dentry_operations {  
    int (*d_revalidate) (struct dentry *, int);  
    int (*d_hash) (struct dentry *, struct qstr *);  
    int (*d_compare) (struct dentry *, struct qstr *, struct qstr *);  
    int (*d_delete) (struct dentry *);  
    void (*d_release) (struct dentry *);  
    void (*d_iput) (struct dentry *, struct inode *);  
};
```



```
struct file {  
    struct list_head      f_list;          /* list of file objects */  
    struct dentry         *f_dentry;        /* associated dentry object */  
    struct vfsmount       *f_vfsmnt;        /* associated mounted fs */  
    struct file_operations *f_op;          /* file operations table */  
    atomic_t               f_count;         /* file object's usage count */  
    unsigned int           f_flags;         /* flags specified on open */  
    mode_t                f_mode;          /* file access mode */  
    loff_t                f_pos;           /* file offset (file pointer) */  
    struct fown_struct    f_owner;         /* owner data for signals */  
    unsigned int           f_uid;           /* user's UID */  
    unsigned int           f_gid;           /* user's GID */  
    int                   f_error;         /* error code */  
    struct file_ra_state  f_ra;            /* read-ahead state */  
    unsigned long          f_version;       /* version number */  
    void                  *f_security;      /* security module */  
    void                  *private_data;    /* tty driver hook */  
    struct list_head       f_ep_links;      /* list of eventpoll links */  
    spinlock_t             f_ep_lock;        /* eventpoll lock */  
    struct address_space   *f_mapping;       /* page cache mapping */  
};
```

```
struct file_operations {  
    struct module *owner;  
    loff_t (*llseek) (struct file *, loff_t, int);  
    ssize_t (*read) (struct file *, char *, size_t, loff_t *);  
    ssize_t (*aio_read) (struct kiocb *, char *, size_t, loff_t *);  
    ssize_t (*write) (struct file *, const char *, size_t, loff_t *);  
    ssize_t (*aio_write) (struct kiocb *, const char *, size_t, loff_t *);  
    int (*readdir) (struct file *, void *, filldir_t);  
    unsigned int (*poll) (struct file *, struct poll_table_struct *);  
    int (*ioctl) (struct inode *, struct file *, unsigned int, unsigned long);  
    int (*mmap) (struct file *, struct vm_area_struct *);  
    int (*open) (struct inode *, struct file *);  
    int (*flush) (struct file *);  
    int (*release) (struct inode *, struct file *);  
    int (*fsync) (struct file *, struct dentry *, int);  
    int (*aio_fsync) (struct kiocb *, int);  
    int (*fasync) (int, struct file *, int);  
    int (*lock) (struct file *, int, struct file_lock *);  
    ssize_t (*readv) (struct file *, const struct iovec *,  
                     unsigned long, loff_t *);  
    ssize_t (*writev) (struct file *, const struct iovec *,  
                     unsigned long, loff_t *);  
    ssize_t (*sendfile) (struct file *, loff_t *, size_t,  
                        read_actor_t, void *);  
    ssize_t (*sendpage) (struct file *, struct page *, int,  
                        size_t, loff_t *, int);  
    unsigned long (*get_unmapped_area) (struct file *, unsigned long,  
                                      unsigned long, unsigned long,  
                                      unsigned long);  
    int (*check_flags) (int flags);  
    int (*dir_notify) (struct file *filp, unsigned long arg);  
    int (*flock) (struct file *filp, int cmd, struct file_lock *fl);  
};
```



VFS structures

```
struct file_system_type {  
    const char          *name;      /* filesystem's name */  
    struct subsystem    subsys;    /* sysfs subsystem object */  
    int                 fs_flags;   /* filesystem type flags */  
  
    /* the following is used to read the superblock off the disk */  
    struct super_block  *(*get_sb) (struct file_system_type *, int,  
                                    char *, void *);  
  
    /* the following is used to terminate access to the superblock */  
    void                (*kill_sb) (struct super_block *);  
  
    struct module        *owner;    /* module owning the filesystem */  
    struct file_system_type *next;   /* next file_system_type in list */  
    struct list_head      fs_supers; /* list of superblock objects */  
};
```





```
struct files_struct {  
    atomic_t    count;           /* structure's usage count */  
    spinlock_t  file_lock;      /* lock protecting this structure */  
    int         max_fds;        /* maximum number of file objects */  
    int         max_fdset;      /* maximum number of file descriptors */  
    int         next_fd;        /* next file descriptor number */  
    struct file **fd;          /* array of all file objects */  
    fd_set     *close_on_exec;  /* file descriptors to close on exec() */  
    fd_set     *open_fds;        /* pointer to open file descriptors */  
    fd_set     close_on_exec_init; /* initial files to close on exec() */  
    fd_set     open_fds_init;   /* initial set of file descriptors */  
    struct file *fd_array[NR_OPEN_DEFAULT]; /* default array of file objects */  
};
```

```
struct fs_struct {  
    atomic_t    count;           /* structure usage count */  
    rwlock_t    lock;            /* lock protecting structure */  
    int         umask;           /* default file permissions */  
    struct dentry *root;         /* dentry of the root directory */  
    struct dentry *pwd;          /* dentry of the current directory */  
    struct dentry *altroot;      /* dentry of the alternative root */  
    struct vfsmount *rootmnt;   /* mount object of the root directory */  
    struct vfsmount *pwdmnt;    /* mount object of the current directory */  
    struct vfsmount *altrootmnt; /* mount object of the alternative root */  
};
```



open() system call

fd = open("/home/sunbeam/abc.txt", O_RDONLY);



open() system call

①
`fd = open("/home/sunbeam/abc.txt", O_RDONLY);`





read() system call

VFS
(logical file
sys)

FS mgr
(Physical file
sys)



IO Subsystem

Driver



inode/FCB

file Blk	Disk Blk
0	100
1	105
2	102
3	107



read(fd, buf, sizeof(buf))

100

sys-read(fd, buffer, length)

vfs-read(file, buffer, length, inode)

ext3_read(file, inode, file_block)

disk_read(disk-device, disk-block)

Buffer
cache



write() system call

VFS
(logical file sys)

FS mgr
(Physical file sys)

IO Subsystem

Driver



lseek() system call

`lseek(fd, +10, SEEK_SET)` $f_pos = 10$

`lseek(fd, -10, SEEK_END)` $f_pos = 9990$

`lseek(fd, -10, SEEK_CUR)` $f_pos = 4990$

`lseek(fd, +10, SEEK_CUR)` $f_pos = 5010$



4990
CUR
5000
5010

9990 10000



Thank you!!!

Devendra Dhande

devendra.dhande@sunbeaminfo.com