

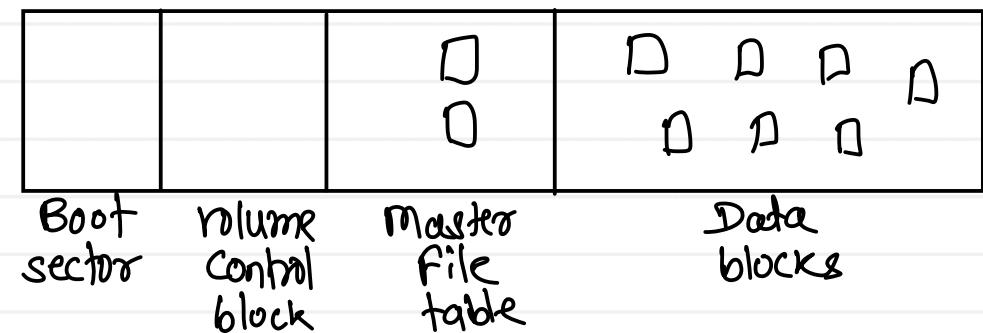
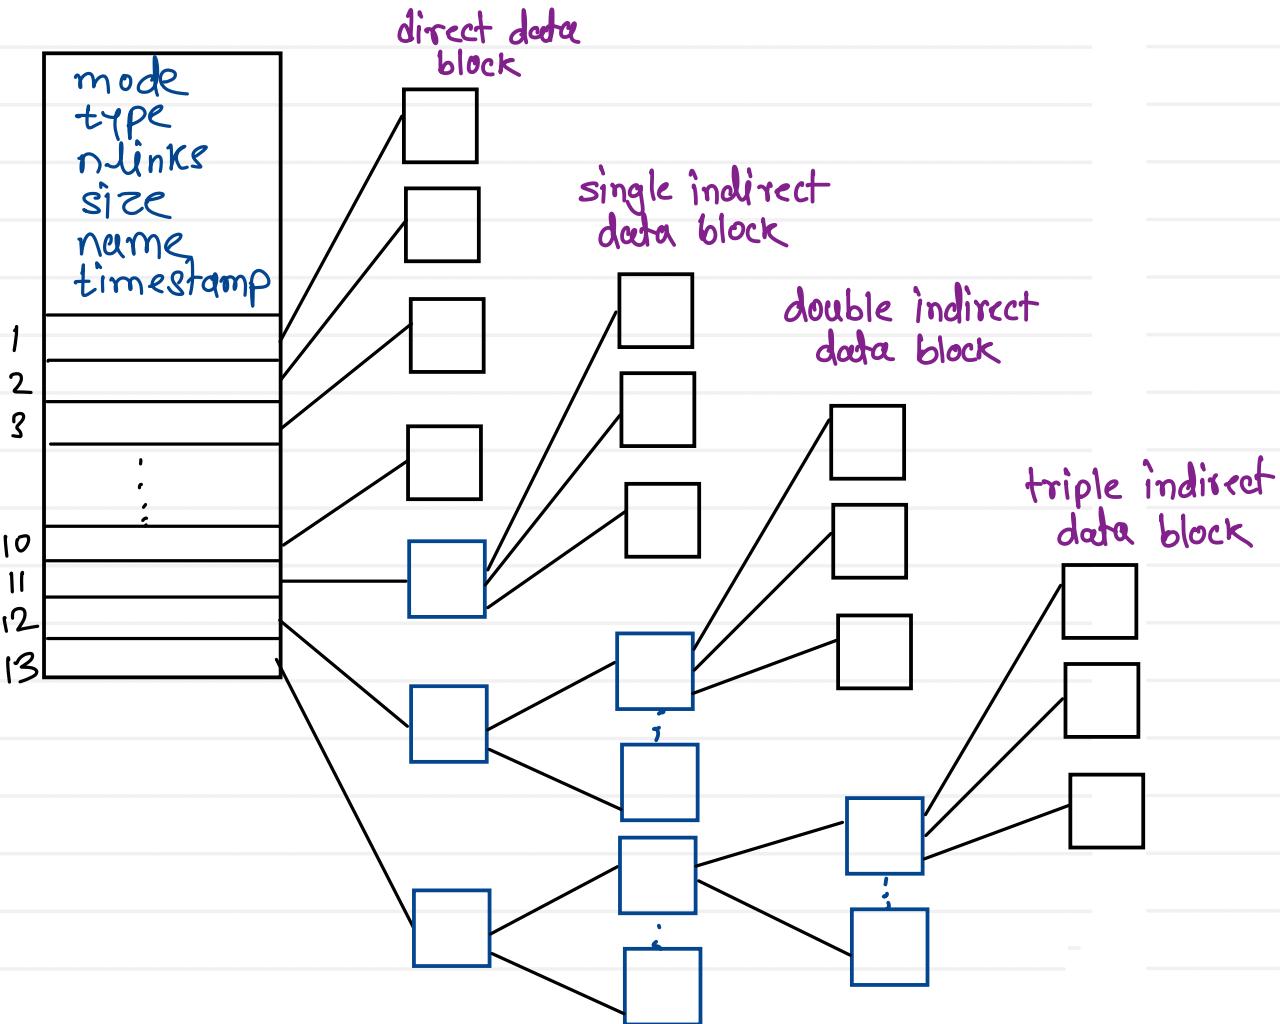


**Sunbeam Institute of Information Technology  
Pune and Karad**

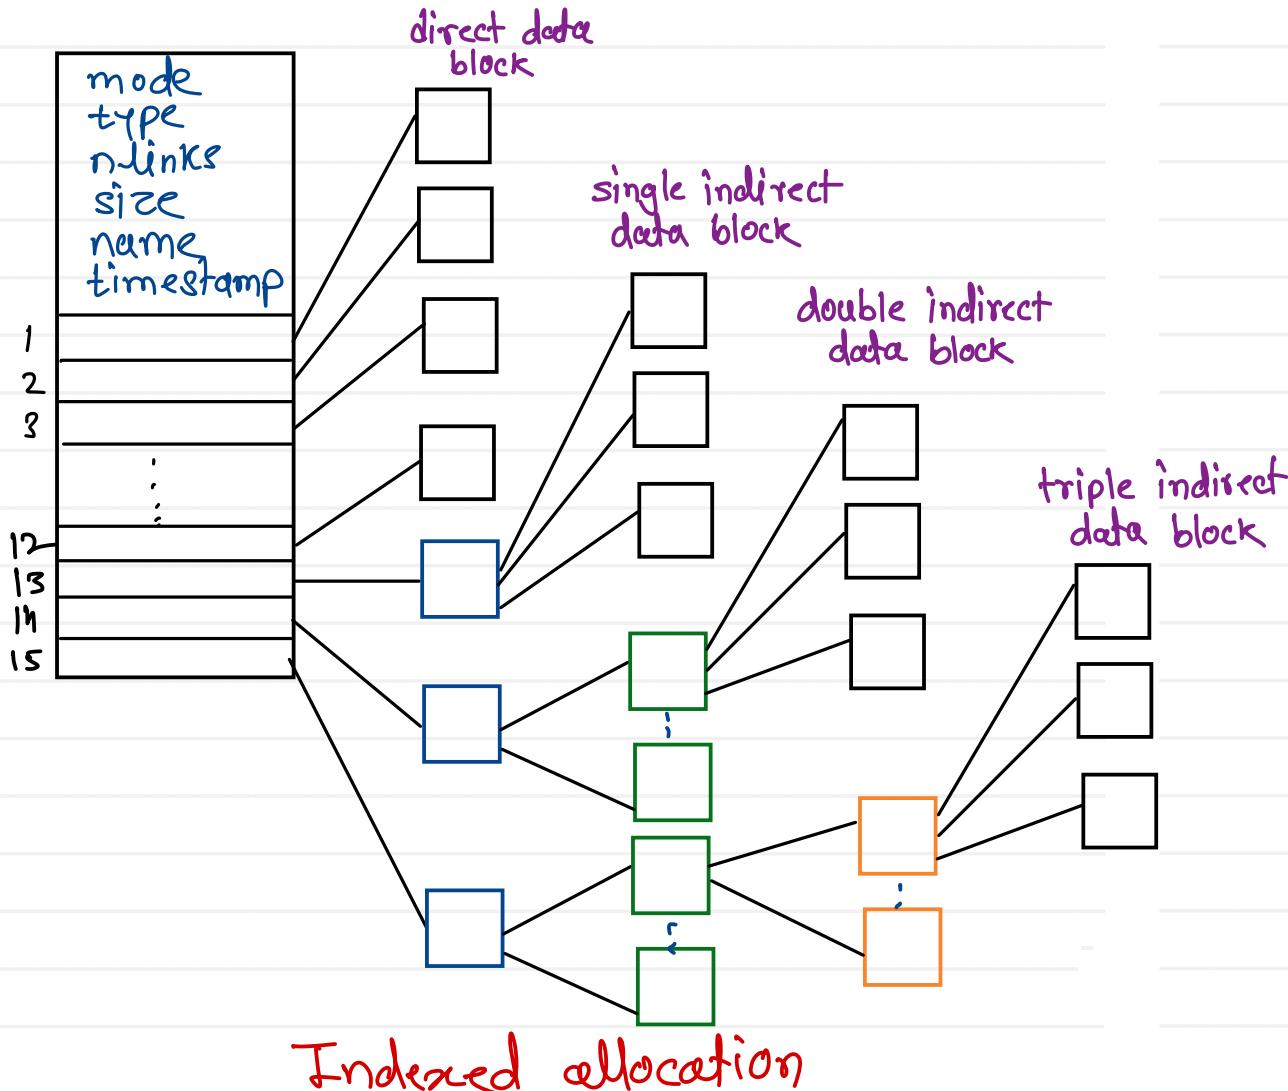
**Module - Embedded Operating System**

Trainer - Devendra Dhande  
Email – [devendra.dhande@sunbeaminfo.com](mailto:devendra.dhande@sunbeaminfo.com)

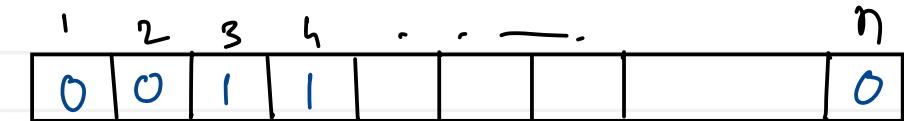
# UFS (UNIX File System)



# Ext file system

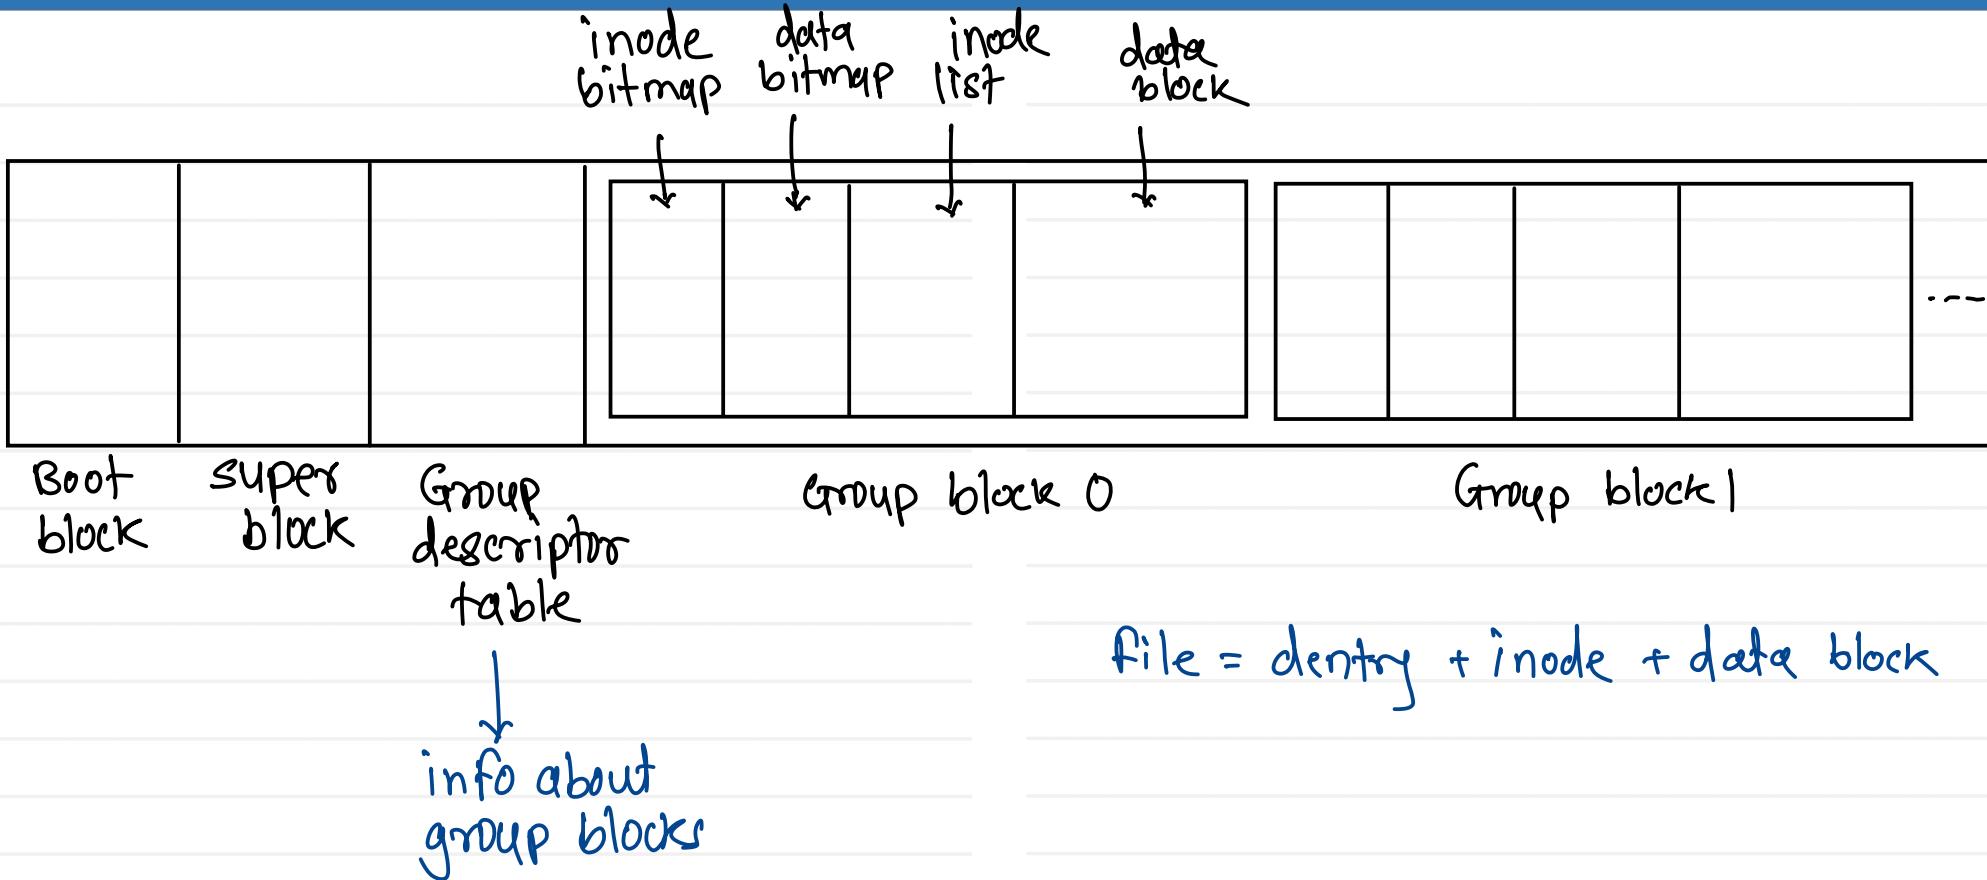


Bitmap

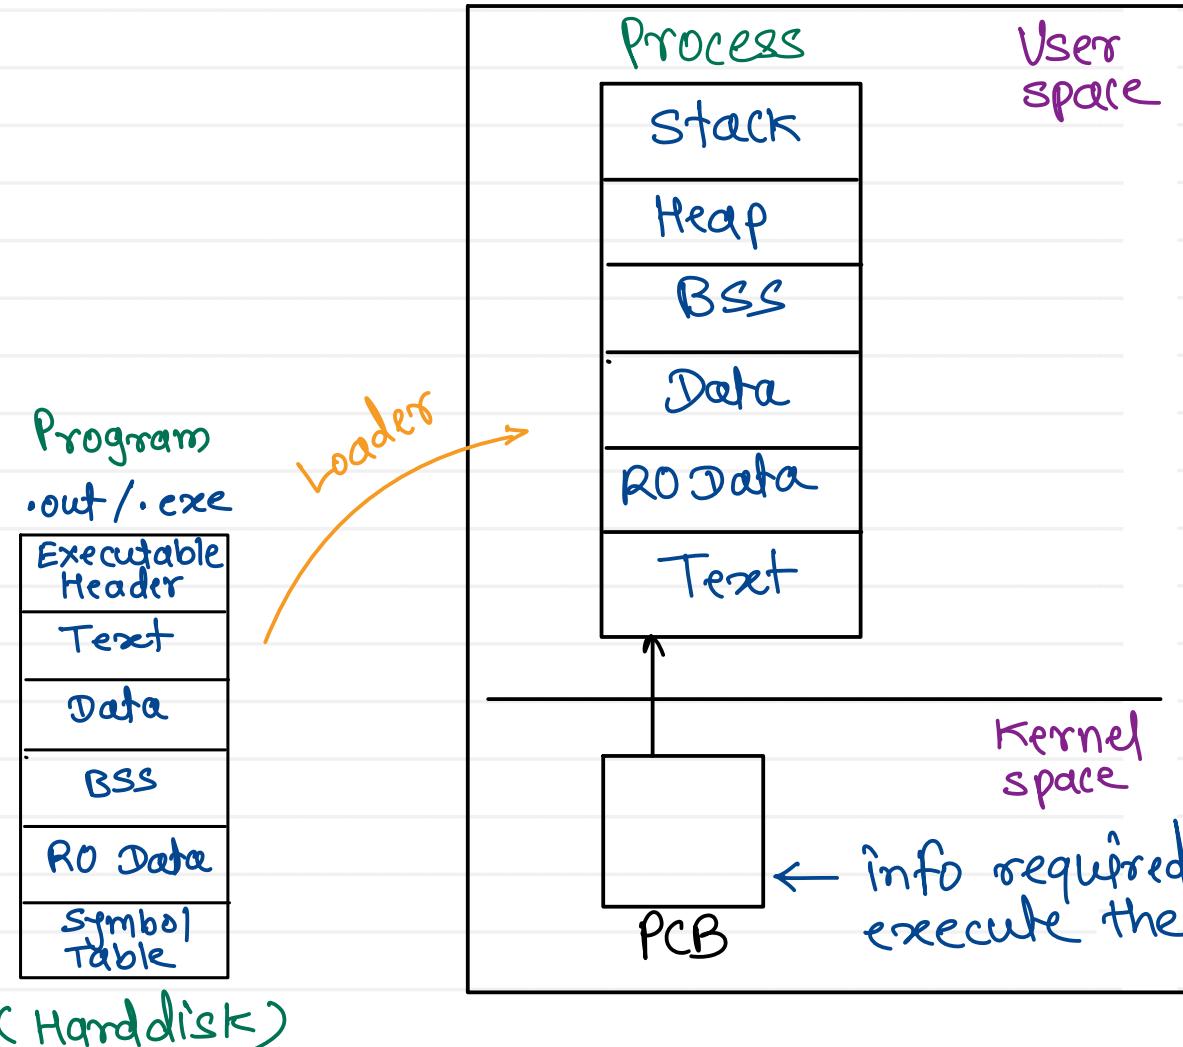


free space management

# Ext Filesystem



# Process



PCB - Process Control Block

↳ Process Descriptor

↓  
struct task\_struct  
< sched.h >

UNIX - uname

- PCB :
1. pid, ppid
  2. exit status
  3. sched info ( pri, state, algo ... )
  4. mem info ( base & limit, page/segment table )
  5. file info ( opened file, cwd ... )
  6. IPC info ( signal, buffers ... )
  7. execution context
  8. kernel stack

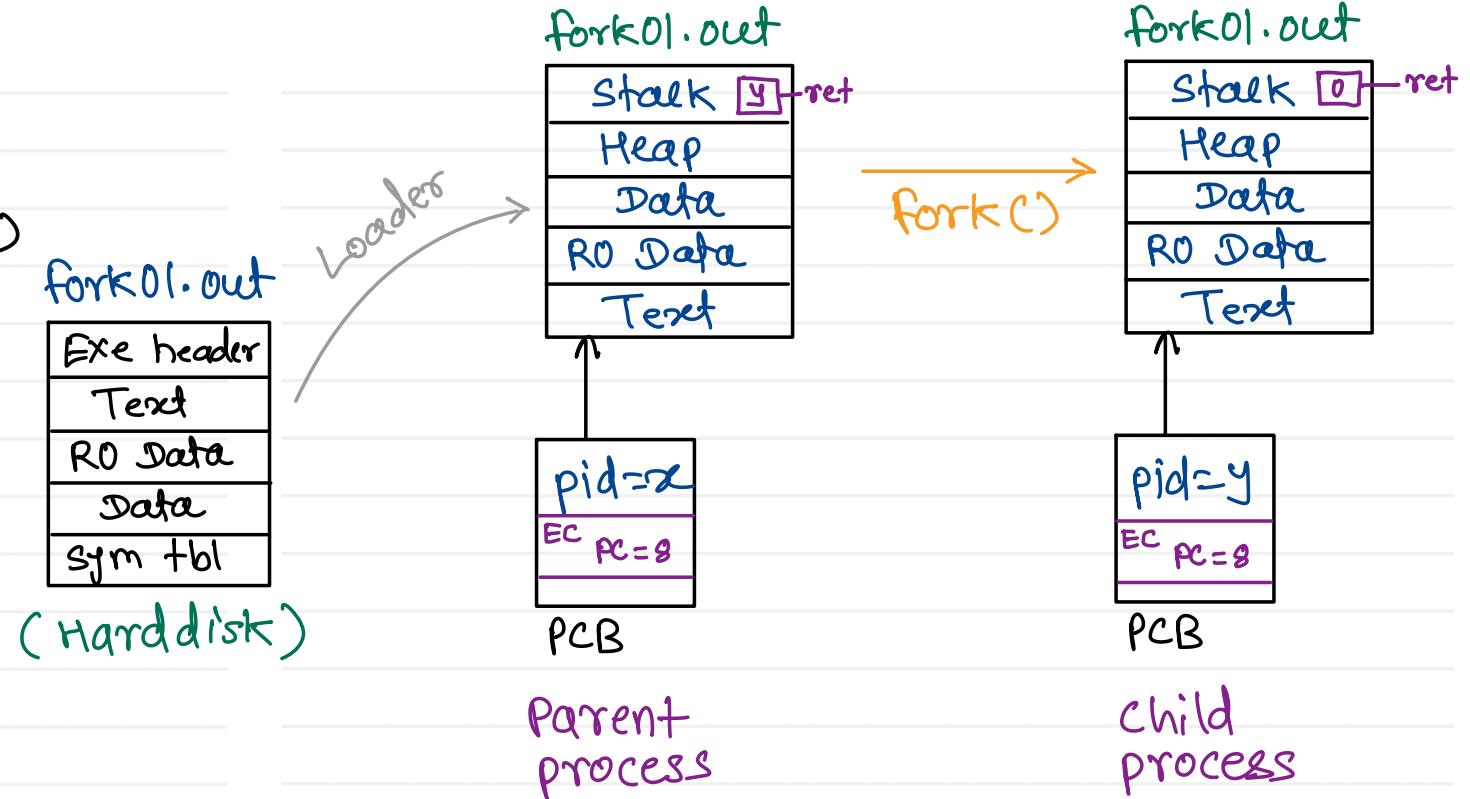


# Process creation

1. Windows : create\\_process()
2. UNIX : fork()
3. BSD UNIX : fork(), vfork()
4. Linux : fork(), vfork(), clone()

```
#include<stdio.h>
#include<unistd.h>

int main(void) {
    printf("program started");
    fork();
    printf("program finished");
    return 0;
}
```



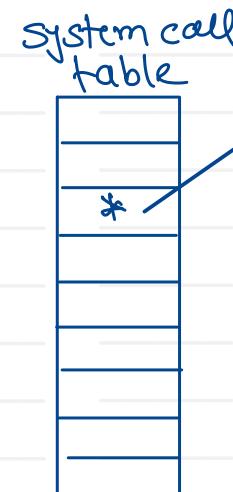
- child process is created by duplicating parent process.
- calling process = parent process,  
new process = child process
- both processes have separate memory spaces.
- both processes are scheduled separately.



# fork() internals

```
int main(void) {  
    int ret = fork();  
    printf("ret = %d", ret);  
    return 0;  
}
```

```
swi_handler() {  
    1. save execution context of current  
       process into kernel stack of process.  
    2. find address of actual system call  
       implementation from system call table  
    3. call system call implementation  
    4. pid = CPU_scheduler();  
    5. restore the execution of selected process  
       from kernel stack on CPU.  
}
```



Return value of fork :

On success :

PID of child is returned into parent  
0 is returned into child

On failure :

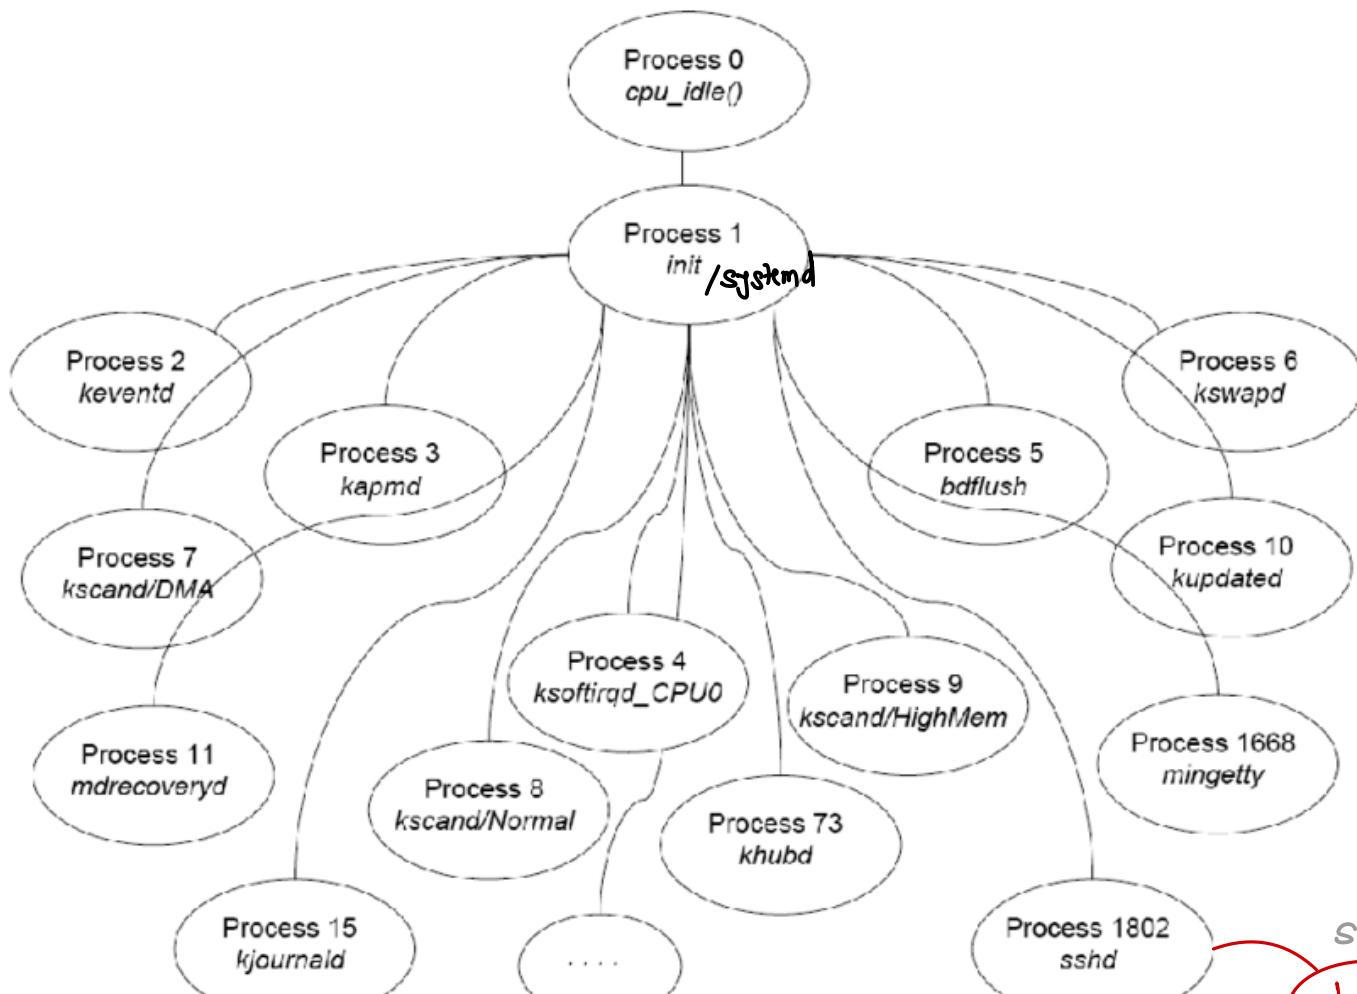
-1 is returned into parent & child  
process is not created.

sys\_fork() {

1. allocate space for PCB of new process
2. copy PCB of parent into child's PCB
3. change few values from child's PCB (pid..)
4. allocate memory in user space
5. copy parent's data into child's space
6. Update return value in both
  - i. parent process - execution context  
 $r0 = \text{pid of child}$
  - ii. child process - execution context  
 $r0 = 0$

}

# Process hierarchy

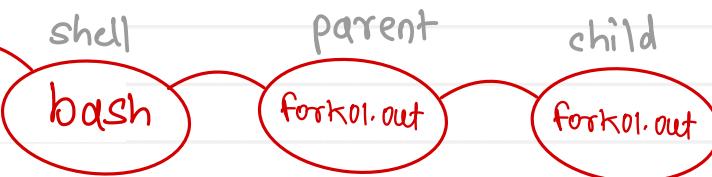


"`init` is first process created using `fork()`. Its pid is always 1.

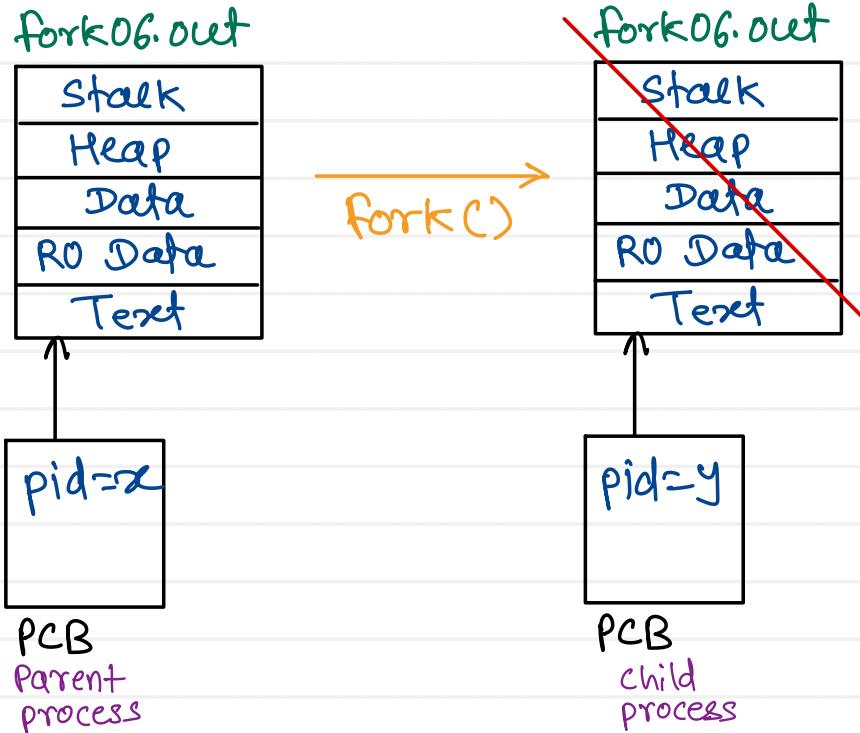
In modern Linux kernels(3.0+), `init` is renamed to `systemd` and kernel initialization/ startup is improved.

`init` was designed for uniprocessor systems, so service startup was sequential.

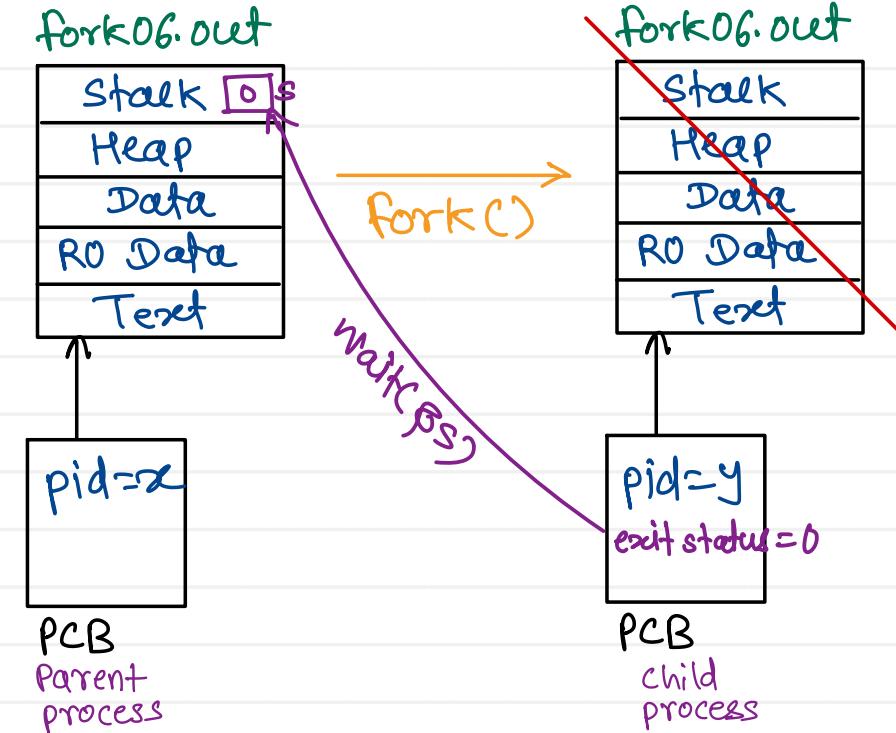
`systemd` designed for multiprocessor systems, so that multiple services can start simultaneously (can work on diff. CPUs). This speed up the system booting.



# Zombie state



- When child is terminated before its parent & parent do not read exit status of child, child becomes zombie.
- PCB of zombie is kept into memory until its parent read its exit status



- Wait will block the execution of parent, will read exit status of child from its PCB & will store it into s variable. State of child will be changed to "DEAD".
- PCB of DEAD process is cleared immediately from RAM.



Thank you!!!

Devendra Dhande

[devendra.dhande@sunbeaminfo.com](mailto:devendra.dhande@sunbeaminfo.com)